Partner, Don’t Police: AI in the Business Classroom

Michael Borck

2026-11-02

Partner, Don't Police
AI in the Business Classroom

Copyright © 2025 Michael Borck.

Book content: Creative Commons Attribution 4.0 International (CC BY 4.0 International).
https://creativecommons.org/licenses/by/4.0/

Original code examples and accompanying software: MIT License.
https://opensource.org/licenses/MIT

These grants cover material the author has authority to license. Separately credited third-party material retains its stated terms.

Published by Michael Borck
Perth, Western Australia

ISBN: 979-8-2542-3619-1

First edition, 2025.

AI disclosure: This book was written using the methodology it describes. AI tools were used as thinking partners throughout the drafting, iterating, and refining process. The author reviewed, challenged, and took responsibility for every sentence.

Companion website:
https://michael-borck.github.io/partner-dont-police
Source: https://github.com/michael-borck/partner-dont-police

Table of contents

Preface 1
Acknowledgments 5
1 Introduction: A Defensible Classroom Pilot 7
2 Understanding AI and Large Language Models . . . 11
3 Ethics, Data Governance & Integrity 17
4 Getting Started: Your First AI Conversation 25
5 Seven Essential Prompt Techniques for Business Teaching 31
6 Managing Context: A Checked Handoff 43
7 The Critique Toolkit: Check the Recommendation . . 49
8 The Flight Simulator: Bounded Conversation Practice 55
9 Self-Assessment: Judge Before Receiving Feedback . 61
10 Assessing Learning: Process, Product and Independent Capability 67
11 Transforming Content: Keep the Learning Decision Visible 75
12 Virtual Company: Keep a Shared Record 79
13 Designing an AI-Supported Unit 87
14 Adapting to a Different Teaching Context 95
15 Implementation Practicalities 101
16 Conclusion: Make the Next Decision 107

Appendices **111**
A Educator Pilot Pack: PDP-2026.1 111
B Prompt Library 123
C Workshop Guide 133
D Aligning the Activity with Institutional Outcomes . . . 141
E A Calibrated Assessment Rubric 145
F Stress-Test the Assessment 151
G Glossary 155
H Further Reading 161
About the Author 165

Preface

This book is for business educators deciding where AI-supported practice belongs in a course, and what students must still demonstrate independently. It offers a small, inspectable starting point: a fictional case, source records, contrasting responses, flawed feedback, a checked revision and a changed task.

You do not need a paid AI account to use the book. The worked examples and activities include paper or paired routes. An institution may reasonably decide that a particular tool, dataset or activity is unsuitable.

Reading routes

The sixteen chapters follow four stages:

- **Foundations (1–3):** the educational purpose, how language models and products differ, and the ethics and data gate to complete before using a service.
- **Practice (4–9):** getting started, seven prompting patterns, managing context, critique, bounded simulation and self-assessment.
- **Design (10–13):** assessment evidence, a small content transformation, the InnovateCo case and a unit-level design.
- **Implementation (14–16):** local adaptation, practical constraints and a decision about a first pilot.

For a first classroom trial, read Ethics and Integrity, then the Pilot Pack and Flight Simulator. For assessment redesign, start with

Assessment, the Rubric System and the Stress Test.

The appendices also provide the full prompt library, a facilitator's workshop guide, institutional alignment, terminology and further reading. They are part of the downloadable editions, not web-only extras.

What the examples establish

PDP-2026.1 is an authored fictional teaching pack, not a report of a classroom trial. Its scores illustrate educator calibration against stated criteria. They do not validate the rubric for another institution or demonstrate improved learning.

The main case concerns a disputed performance issue. It teaches evidence handling and a bounded first response, not qualification to conduct a real investigation. Supplier and marketing examples show how the same checks work outside HR. Adapt the outcome and sources before adopting an example in your discipline.

The book distinguishes supported practice, submitted process evidence and independently demonstrated capability. A polished AI-assisted response, a long transcript and an AI-generated mark cannot substitute for the last of these.

Formats and supporting resources

The online edition (https://michael-borck.github.io/partner-dont-police) is the primary, freely accessible edition. The site provides PDF, EPUB and plain-text downloads, plus the source and related resources.

Paperback ISBN: 979-8-2542-3619-1. Readers who want to support the open edition can look for the paperback on Amazon (https://www.amazon.com.au/s?k=9798254236191). Availability and the edition supplied should be checked on the listing; the ISBN is not evidence that a print copy contains the latest source revisions.

The GitHub repository (https://github.com/michael-borck/partner-dont-police) contains the manuscript and history. DeepWiki (https://deepwiki.com/michael-borck/partner-dont-police) offers an AI-generated view of that repository; check its answers against the source. Any companion chatbot is similarly a navigation aid, not an authoritative interpreter of policy or assessment requirements. Do not submit confidential material to it.

The downloadable `llm.txt` can support local search or an approved AI workflow. Its length may exceed a product's usable context, and uploading it does not ensure every passage will be considered.

Scope and currency

This revision was prepared in September 2026. Product capabilities, institutional rules and legal requirements can change; check the relevant current authority before implementation. Australian examples require local adaptation elsewhere. Nothing here replaces institution-specific approval or professional advice about a real case.

Conversation, Not Delegation explains the wider methodology; this book includes the recap needed to work independently. Find the companion and other open books at books.borck.education (https://books.borck.education).

Acknowledgments

The ideas in this book were developed in classrooms, not in the abstract. The students in my business education units were the first to encounter these approaches (the flight simulators, the critique exercises, the process-based assessments) and their responses shaped every chapter. Some approaches worked immediately. Others failed in ways that were far more instructive than the successes. Both made the book better.

Colleagues across business disciplines provided feedback, scepticism, and the occasional "have you tried this with accounting students?" that sent entire chapters in new directions. Teaching and learning teams provided the institutional context that kept the ideas grounded in what is actually possible within a university semester.

This book draws on the methodology developed in *Conversation, Not Delegation*, and the two books were written in parallel, each sharpening the other. The pedagogical frameworks it builds on (experiential learning, reflective practice, authentic assessment) are not new. What is new is applying them through AI, and that application was only possible because decades of educational research provided the foundation.

The technical infrastructure behind this book is entirely open source. Quarto, Python, GitHub, and GitHub Pages made it possible to write, build, and publish across multiple formats without a traditional publisher. The open source community deserves more credit than it typically receives.

AI tools were used throughout the writing process, as described

in the preface of the companion book. Claude (Anthropic) served as a conversation partner for drafting, iterating, and refining. The process was exactly what both books advocate: conversation, not delegation. The author's judgement shaped every page. The AI made the work faster. It did not make the decisions.

Chapter 1

Introduction: A Defensible Classroom Pilot

Start with a learning decision you can examine.

1.1 The problem this book addresses

A business educator may want students to practise difficult decisions more often than a timetable allows: asking a fair question, checking a recommendation, explaining an uncertainty or revising an unsupported conclusion. AI can help generate practice and feedback, but it can also make weak reasoning look fluent.

This book offers a way to design and evaluate that practice. The aim is to help you decide where assistance is useful, where checks are needed and how the learner will demonstrate capability. More AI use is not by itself the measure of success.

You bring disciplinary and teaching judgement. A model may contribute alternative wording, questions or a simulated role. Neither its confidence nor the length of the resulting conversation establishes what a student knows.

1.2 A small example of the difference

In the fictional pilot, a junior adviser meets an employee whose manager has proposed a performance plan. A polished response announces the plan, invents an earlier warning and offers to approve support arrangements. It sounds empathetic, but conflicts with the supplied records and authority.

A shorter response checks the missing evidence, preserves the opportunity to respond and refers the decision appropriately. It leaves the matter unresolved. That is the stronger first response because its reasoning fits the facts.

The pilot pack provides both responses, flawed feedback, a correction and a changed case. You do not have to invent a scenario or trust a live model to work through the book.

1.3 The reader and the outcome

The main reader is an educator teaching business or a related professional discipline. You do not need to code or purchase an AI subscription. You do need enough disciplinary knowledge to recognise the task's standards and arrange expert review where the material exceeds your competence.

Student-facing prompts and instructions are labelled as such. They are examples to adapt, not institution-wide policies or ready-made legal advice.

By the end, you should be able to:

- Specify a bounded capability and the evidence needed to assess it.
- Select permitted data and an approved or non-AI route.
- Design practice and feedback with checkable inputs.
- Inspect a response and justify a correction or warranted agreement.
- Include an appropriate independent demonstration.
- Evaluate access, workload and learning evidence before expanding a pilot.

A model's success at producing a handout is not the same as your success at designing an assessment. The latter requires these decisions to fit together.

1.4 Conversation as a working method

The companion *Conversation, Not Delegation* develops the broader method of thinking with AI. PDP uses a short practical version: clarify the question, consider alternatives, check consequential claims and make a decision you can defend.

Its VET reminder is useful here: **Verify** the claim, **Explain** your reasoning and **Test** it against a changed condition. The classroom application adds a further question: which part of the student's performance was assisted, and which part demonstrates independent capability?

A conversation can support that work. It can also become a long rehearsal of unsupported assertions. Do not reward the conversation itself without examining the decisions it contains.

1.5 Three design commitments

First, establish safety before use. The data gate covers permission, account features, handling and alternatives. Fictional material is the default for practice; a real placement case is not automatically suitable because names were removed.

Second, make feedback contestable. A model can misapply a rubric or invent a fact. Students need supplied criteria and examples that let them check its comments. They should be allowed to retain correct advice as well as reject errors.

Third, preserve the learner's opportunity to act independently. The final changed case cannot be replaced by a reflection that merely says learning occurred. Choose a demonstration appropriate to the outcome and access needs.

1.6 The route through the book

Begin with Understanding AI and Ethics and Data Governance. Then try Getting Started, choose a prompt pattern and practise a checked handoff.

The next sequence develops critique, simulation, self-assessment and assessment design. Small content activities and the virtual company then lead into whole-unit design.

The closing chapters address adaptation, implementation and a pilot decision. The appendices supply the cases, rubric, stress test and workshop resources. These support a standalone route; the other books are optional extensions.

1.7 Activity: choose one capability

Allow eight minutes; no AI required. Pick a task you teach. Complete:

> After this activity, the student can ___ using ___ evidence, under ___ conditions.

Avoid "understand AI better" or "produce a professional report" without saying what would demonstrate that outcome. In the pilot, the target is to plan and justify a fair first response using supplied records. It does not claim to certify a competent workplace investigator.

Keep your sentence. Each later chapter should help you make a concrete design decision about that capability.

Chapter 2

Understanding AI and Large Language Models

> A useful explanation of a tool should help you decide how to check its work.

2.1 Model, product and task

A language model generates text from patterns learned during training and the context supplied at use time. Many contemporary models are trained to predict tokens, which may represent words or parts of words, and undergo additional training to improve their behaviour. This mechanism supports useful performance on many tasks; it does not guarantee factual accuracy.

A product may wrap a model with search, document retrieval, calculation, code execution or memory features. Those additions matter. A reply based on a retrieved document has a different evidence trail from an unsupported answer generated from learned patterns, but either can be wrong.

ChatGPT, Claude, Gemini and Copilot are product families, not interchangeable specifications. Features, accounts and institutional approvals change. This book's workflow does not depend on a particular subscription or interface.

The GPT-4 technical report describes next-token pre-training and subsequent fine-tuning, while withholding model size. Claims that its parameter count is known precisely should not be treated as established facts. GPT-4 technical report, section 2 (https://arxiv.org/pdf/2303.08774).

2.2 Why fluent answers need checks

A plausible explanation can contain an invented citation, a wrong calculation or a confident inference unsupported by the inputs. Fluency makes the answer easy to read; it is not a confidence interval or a guarantee.

In the pilot's P1-F feedback, the invented April warning sounds specific. The source card records only a May conversation. Checking the date is a stronger test than asking the same model whether it is certain.

Avoid the opposite overstatement too. "AI cannot reason" is not a useful rule for designing a class. Models can perform tasks involving reasoning, with reliability that varies by task, inputs and system. Observe the actual performance and verify what matters. Neither impressive output nor a blanket dismissal tells you whether the student has learned.

2.3 What a model may help you do

You might ask it to propose alternative classroom questions, play a bounded role, reorganise a draft or suggest explanations at different levels. Each is a candidate use, not an automatic improvement over your current method.

Match the check to the task:

Task	Useful check
Explain a concept	Compare with the taught definition and a counterexample

Task	Useful check
Summarise a record	Check retained facts, omissions and source locations
Calculate a total	Recalculate from the inputs
Suggest a scenario	Check factual consistency, assumptions and suitability
Give feedback	Compare with a calibrated rubric and specimen responses

A “creative” task can still cause harm through stereotypes, personal data or false claims. A small task can be consequential. Consider the stakes and whether you can detect an error, not only the size or apparent precision of the request.

2.4 Context, retrieval and memory

The active context is the information made available for the current response. It has limits, and a product may truncate, summarise or select parts of a conversation. Even when information fits, the model may not use it reliably.

Some products can retain preferences or retrieve earlier material; others do not. Do not assume either perfect continuity or universal amnesia. For a teaching simulation, the authoritative record should be the source cards and event log, not what a persona seems to remember.

Long-context research found that relevant information’s position can affect performance, sometimes with stronger results near the beginning and end than the middle. That is a warning to test retrieval, not a universal 50-turn rule. Lost in the Middle (https://arxiv.org/abs/2307.03172).

The context chapter shows how to carry a checked brief between sessions. Do not copy unrelated private material just to preserve context.

2.5 Generated perspectives are not independent sources

Three personas can help you imagine competing concerns. They do not represent three people consulted, three experiments or three independent confirmations. They may repeat the same error because they draw on the same inputs and model.

In a business class, ask students to distinguish a possible stakeholder concern from evidence that stakeholders actually hold it. A model's "employee" can suggest a question to investigate, but cannot establish an employee's experience.

Likewise, an explanation of how a model reached an answer is not an inspection of its internal computation. Ask for checkable steps, sources, assumptions and calculations. Evaluate those directly.

2.6 The learner remains a separate question

A good output and a learning gain are different outcomes. Assisted practice can improve the work in front of the student without improving their later unaided performance. The design of assistance matters.

One study of high-school mathematics compared ordinary and guardrailed AI support with unaided follow-up performance. It is evidence for checking the difference between assisted performance and learning, not a validation of every AI activity in this book. Bastani et al., author manuscript (https://hamsabastani.github.io/education_llm.pdf).

In PDP, the learner therefore attempts a response, examines feedback and then tries a changed case independently. The model is optional in that sequence. The method requires evidence of the learner's capability, not loyalty to a tool.

2.7 Activity: separate claim from check

Allow eight minutes; no AI required. Read the P1-F feedback in the pilot pack. Underline one factual claim, one judgement and one recommendation. Name the source or criterion needed to check each.

A successful answer checks the April warning against P1-D3, the fairness judgement against P1-D1 and the approval recommendation against P1-D5. "Ask the AI again" may generate another explanation, but is not independent verification.

Chapter 3

Ethics, Data Governance & Integrity

> Decide what may enter the tool before asking what the tool can do with it.

3.1 A decision before the first prompt

This chapter belongs before the first activity because a useful output does not justify an inappropriate input. Use the fictional pilot pack while learning. Do not substitute a real student, employee, client or research participant to make the exercise more realistic.

Apply the same first-use gate to drafting, simulations, feedback, research and workshops:

1. **Classify the material.** Is it public, confidential, personal, sensitive, contract-restricted or otherwise controlled? Who might be identifiable from combinations of details?
2. **Check authority and rights.** Are you permitted to process it for this purpose and share it with this service? Possessing a document is not the same as having permission to upload it.
3. **Check the approved account and features.** Institutional approval should cover the use, account, data class and enabled features, not just a brand name.

4. **Check handling.** Consider retention, training use, access, sharing, external search/integrations and relevant processing locations. Ask the institution's appropriate adviser where information is unclear.
5. **Choose a route.** If the requirements are not met, use synthetic material, an approved alternative or the non-AI version. Do not upload first and seek permission afterwards.

Record the decision briefly, without copying the sensitive information into the record. An approved service may still be inappropriate for a particular dataset or purpose.

3.1.1 What enterprise protection does and does not establish

Enterprise services can offer contractual protections and administrative controls. They do not automatically establish your institution's permission for every use. Features and connected services can have different boundaries. Microsoft, for example, distinguishes Copilot protections from web-search processing. Check the actual service and configuration. Microsoft privacy and protections (https://learn.microsoft.com/en-us/copilot/privacy-and-protections).

OAIC recommends due diligence appropriate to an AI product's intended use, including access, oversight and privacy/security risks. Its guidance also recognises re-identification risk. Removing names is not a universal clearance to use an unapproved tool. OAIC guidance on commercial AI products (https://www.oaic.gov.au/privacy/privacy-guidance-for-organisations-and-government-agencies/guidance-on-privacy-and-the-use-of-commercially-available-ai-products).

De-identification is a contextual risk decision, not a "find and replace names" exercise. A rare role, incident date, small cohort or distinctive quotation may still identify someone. It also does not remove copyright, contractual or research-ethics restrictions. For classroom practice, a genuinely invented case is usually easier to govern than a disguised real case.

3.2 Keep different data risks separate

A conversation can be retained as a service record without being used for model training. If it enters training, the resulting model is not simply a searchable folder of every document; nevertheless, memorisation and extraction can occur. A retrieval-enabled system can also access stored documents through a separate mechanism. These are different pathways.

Published extraction research shows that some training text can be recovered from language models. That rules out blanket claims of impossibility. It does not show that another user can retrieve any arbitrary one-off upload. Carlini et al. (https://arxiv.org/abs/2012.07805), Nasr et al. (https://arxiv.org/abs/2311.17035).

Training exclusions address one pathway, not every retention, sharing or security risk. Likewise, similar ideas appearing elsewhere establish neither theft nor independent invention. Teach students to identify the evidence needed for a claim instead of substituting reassurance or suspicion.

Credentials, access tokens and unnecessary personal information do not belong in classroom prompts. If sensitive material is accidentally uploaded, stop the activity, preserve appropriate incident details without further disclosure and follow institutional reporting procedures. Do not ask the model to certify that the information has been deleted.

3.3 Transparent use with real boundaries

Students should know what help is permitted, what must be disclosed and what they must demonstrate independently. Permission is task-specific. An institution may legitimately restrict assistance at a particular assurance point while supporting its use during practice.

A suitable statement might say:

```
In this activity, use only the supplied fictional
pack.
You may rehearse with a partner, the printed
exchange
or the approved institutional AI service.

Submit your response and two decision notes.
Explain
what assistance you used and which claims you
checked.
Do not submit unrelated private conversations.

The changed-case task is completed independently
under
the conditions stated in the assessment brief.
Access adjustments are arranged before the task.

Ask if a tool, data source or form of help is
unclear.
Institutional academic-integrity procedures apply.
```

Adapt this to the actual unit policy. Do not claim that disclosure makes every use permissible, or that any uncritical use automatically meets the institution's definition of misconduct. Poor work and dishonest conduct are different questions.

3.4 Detection and assurance

A detector result is not direct evidence of the student's learning process. Neither is a long, polished conversation. A student may generate a plausible reflection; a genuine learner may write briefly or struggle to explain under pressure.

The assessment chapter therefore combines selected process evidence with meaningful independent demonstrations. These support a learning judgement but do not remove the need for fair misconduct procedures.

Avoid giving students mixed messages. If the assignment permits extensive assistance, do not later penalise its mere use. If it restricts assistance, explain the restriction and the assurance purpose. Follow-up questions should be relevant, accessible and consistent with published expectations.

3.5 Teach verification without forced opposition

Critical engagement does not mean always finding fault. A learner should be able to accept a correct calculation, reject an unsupported inference and say that there is not enough information to decide.

The pilot's marketing example illustrates all three. The higher click-through rate can be calculated. The claim of greater profit lacks cost data. A causal claim cannot be established from different audiences and weeks alone. The same response can reasonably contain agreement, rejection and uncertainty.

A model-generated persona is not an independent expert. Agreement between three personas does not create three sources. Check against the supplied records, reliable external material or an authorised person, depending on the question.

3.6 Sensitive scenarios and student wellbeing

Performance plans, discrimination complaints and termination conversations can be useful teaching contexts, but students may have personal experience of them. Provide advance notice, a stop instruction, a neutral route and a debrief.

Do not reward disclosure of private information, emotional conformity or success at making the simulated employee agree. A respectful decision may be to pause, seek independent advice or leave the outcome unresolved.

For the pilot, the invented policy is the exercise criterion. It is not legal advice. If teaching actual Australian workplace bullying law,

use the complete applicable definition and scope: health/safety risk and the distinction for reasonable management action matter, as well as repeated unreasonable behaviour. Fair Work Ombudsman guidance (https://www.fairwork.gov.au/employment-conditions/bullying-and-harassment). Have an appropriately qualified person review consequential legal teaching.

3.7 Equity is a design requirement

Conversational skill does not erase differences in access, language support, model capability, disability or available time. Provide a route that uses the same learning criteria without a paid subscription or personal account.

When the capability is planning and justification, paper or an accommodated spoken response may work. When the capability is live communication, a written case analysis alone is not an equivalent demonstration: provide an appropriate interaction, such as a structured peer/facilitator role-play.

Do not require every private chat as proof of effort. Ask for bounded extracts and permit working notes where appropriate. Set retention and access rules for submitted evidence through institutional processes.

3.8 A bounded ethics discussion

A fictional recruitment tool ranks applicants with career gaps lower. Three candidates question the decision. The supplied facts do not establish the cause or legal outcome.

Ask learners what evidence would distinguish a relevant job criterion from an unfair proxy, who is authorised to review the system, and what should happen while the concern is investigated. They should not diagnose the algorithm by asking another model whether it is biased.

Speculative systemic risks can be worthwhile extension topics, but

should be labelled scenarios rather than predictions. A claimed "95% accurate" system needs a defined population, outcome, error distribution and validation method before that number can inform an ethical decision.

3.9 Activity: make the upload decision

Allow 10 minutes; no AI required. Apply the five-step gate to:

- The fictional P1 pack.
- A placement client's spreadsheet with names removed.
- A student's draft containing a personal disclosure.

For each, write what you know, what permission or handling detail is missing and the safe alternative. The first can support a bounded practice activity after the account/use check. The second and third need further authority and risk assessment; removing names or choosing an enterprise logo is insufficient.

Keep your resulting rule beside the first student-facing prompt. Safety advice that appears only in a later appendix will not guide the upload decision.

Chapter 4

Getting Started: Your First AI Conversation

> Begin with a task you can inspect and evidence you can check.

4.1 Choose a safe route first

Complete the first-use data gate before opening a tool. Use only the fictional pilot materials or your own non-sensitive invented example. Institutional approval must cover the account, features and intended use; an enterprise label alone is not permission to upload course or student data.

The exercise works in a browser-based approved service if available. Otherwise, compare the supplied responses below on paper. You do not need a paid account to learn how context and criteria change the usefulness of a request.

4.2 Your First Prompt: The Weak Version

If you have an approved tool, try this request. Otherwise use the authored excerpt below:

```
Write a case study about business.
```

The response may be generic, such as a story about a marketing campaign or financial decision. The request supplies little basis for meeting your specific teaching needs. “Business” covers everything. “A case study” could be 100 words or 5,000. The AI had to guess what you wanted.

4.3 Your Second Prompt: A More Specific Version

Now try this instead:

```
You are an expert marketing lecturer at a
university level.

I need you to create a case study for my
undergraduate students that will help them
practice strategic market analysis and
competitive positioning.

Here are the requirements:
- The scenario should involve a product launch
  decision where the company faces competitive
  threats
- Include enough detail that students need to
  identify which market analysis principles apply
- The scenario should be 300-400 words
- End with three discussion questions requiring
  students to analyse the market strategy and
  positioning

Begin.
```

The second request gives a clearer task. It may produce more relevant material, but still needs checking. Here are two **authored illustrative excerpts**, not outputs from a measured comparison:

Generic excerpt: "A business wants to increase sales. It launches a new product. Discuss the marketing strategy."

More specific excerpt: "A small drink producer can fund only one launch channel. Retail buyers want discounts; direct sales need fulfilment capacity. The producer has no reliable repeat-purchase estimate. Which evidence would you request before choosing a channel?"

The second creates a decision and missing information. It still needs enough supplied data and a criterion for evaluating the student's response. Do not mistake a more detailed story for a complete classroom activity.

4.4 The CRAFT Framework

What made that second prompt more specific? It included five elements that you can remember with the acronym **CRAFT**:

Letter	Element	What it does
C	**Context**	Sets the situation: your course, your students, their level
R	**Role**	Defines the AI's perspective: expert lecturer, consultant, practitioner
A	**Action**	Specifies the task: create, analyse, critique, explain
F	**Format**	Describes the structure: word count, number of questions, table format
T	**Tone/Target**	Sets the style and level: professional, introductory, sophisticated

You do not need every element every time. But when a prompt gives you something vague or generic, check which CRAFT elements are

missing. The answer is usually context or format.

CRAFT is not the only prompting framework. The companion book *Conversation, Not Delegation* (available free at books.borck.education) compares several frameworks side by side. Use CRAFT to think intentionally about what you are asking the AI to do, rather than as a rigid formula.

4.5 Conversation, Not One-Shot

CRAFT is your first message, not your only message.

An initial response may miss your teaching context or contain unsupported assumptions. Check it before deciding whether another exchange would help.

The solution is to have a conversation:

```
PROMPT 1 (CRAFT):
"You are an experienced marketing director.
Create a market analysis case study for my
second-year students. 500 words with 4
discussion questions."

[AI produces a case study]

PROMPT 2 (Refine):
"Good. Make the discussion questions more
specific - ask students to identify which
of Porter's Five Forces apply and what data
they would need to support their recommendation."

PROMPT 3 (Expand):
"Add a section comparing three analytical
frameworks and how each leads to different
conclusions."

PROMPT 4 (Repurpose):
"Turn this into a 10-minute classroom activity
```

```
with a 2-minute intro, 5-minute group task,
and 3-minute discussion."
```

You had a conversation in which each follow-up built on what was working, which is what *conversation, not delegation* means in practice.

4.6 Discover Your Strategy: Using AI to Help You Use AI

A model can suggest ways to use a tool, but its advice about current features or policy may be wrong. Check those claims against the product's documentation and your institution's rules.

Instead of reading generic advice, you can have a personalised consultation where AI interviews you about your teaching context and then recommends both obvious and unexpected applications. Copy this prompt and try it now:

```
You're an AI expert consultant for business
educators. I'd love your help discovering
how I could better use AI in my teaching.

Please ask me one question at a time. I'll
answer, then you'll ask the next question.
Ask at most five questions using fictional or
non-sensitive context. Stop if I say STOP.
Focus on:

- My discipline and the students I teach
- My main teaching challenges
- My key teaching goals
- My constraints (time, resources, skills)
- What I've already tried with AI

Once you have enough context, provide TWO
types of recommendations:
```

```
1. OBVIOUS OPPORTUNITIES: Clear, straightforward
   AI applications most educators think of

2. NON-OBVIOUS OPPORTUNITIES: Unexpected or
   creative uses specific to my discipline
   that I might not have considered

Format your recommendations clearly with
implementation tips for each.

Please start by asking your first question.
```

Spend 10-15 minutes answering honestly. The recommendations are only part of the value; you also learn about your own teaching by having to explain it. Writing down those constraints can clarify the task whether you discuss them with a model, a colleague or on paper. Treat the model's recommendations as proposals to evaluate.

> Try this (15 minutes)
>
> Run the consultation prompt above. Answer the questions honestly, and don't polish your responses. Choose one proposal to evaluate against the outcome, data gate, access requirements and staff time. Trying it in class is a separate decision. You can carry out the same planning exercise on paper.

4.7 Activity: inspect the proposed task

Allow ten minutes; no AI required. Compare the two supplied excerpts. Name one decision the learner must make, one missing input and one criterion you would use to assess the response. Then adapt the stronger request to your own discipline without adding real personal or confidential information.

The pilot pack supplies the next step: fixed inputs and worked checks. A usable prompt is only the beginning of a usable learning activity.

Chapter 5

Seven Essential Prompt Techniques for Business Teaching

> The right technique is not the cleverest one. It is the one that matches the kind of thinking you want your students to do.

5.1 Why These Seven?

These seven prompting patterns offer different ways to organise practice. Their educational value depends on the task, evidence, feedback and learner activity; the list is not a claim that each pattern has been validated across every discipline.

1. **Reverse Prompting**: teaches comprehensive scoping and requirement gathering
2. **Pros and Cons**: builds analytical decision-making skills
3. **Stepwise Tutoring**: reinforces process adherence and documentation
4. **Role Play**: develops communication and interpersonal skills
5. **Debating**: strengthens strategic thinking and ethical reasoning

6. **Formative Assessment Generator**: supplies draft practice questions and provisional feedback
7. **The Expert Panel**: develops multi-perspective analysis and synthesis skills

Each pattern needs adaptation to the discipline, supplied records and intended capability. The examples below use different disciplines to show the range, so adapt the prompts for your own field.

These seven are adapted from *Conversation, Not Delegation*, which covers the same approaches for a general audience. If you want the underlying rationale for why each technique works, see the companion book.

5.2 Technique 1: Reverse Prompting (for Scoping)

Instead of solving a problem immediately, the AI asks the student questions to help scope out all requirements, considerations, and potential issues. This mirrors professional work where defining the problem is half the battle.

Example (HR: policy design):

```
I need to draft a new "Flexible Work
Arrangement" policy for a 500-employee company
with both office-based and remote staff.

Your task: Ask me a series of questions to help
me clarify all the requirements, considerations,
and potential pitfalls for this policy.

Ask one question at a time. Wait for my response
before asking the next. Continue until you've
helped me think through at least 10 different
aspects.
```

```
Begin with your first question.
```

The AI asks questions the student must think through before responding: fairness, logistics, legal compliance, technology needs, cultural impact. The questions may broaden their scope, but cannot guarantee completeness. Compare the resulting scope with the supplied task criteria.

Adapt for your discipline: Replace the policy topic with whatever scoping challenge fits: a market entry strategy, an audit plan, a supply chain redesign, an IT implementation. The technique is the same: AI questions, student thinks.

Teaching tip: Use as a pre-writing exercise. Ask for a short evidence index and the subsequent draft. Assess the reasons for consequential decisions; a complete private chat history is not required.

5.3 Technique 2: Pros and Cons (for Decision Making)

The AI systematically analyses multiple approaches to a problem, evaluating each against specific criteria. Students must then critically evaluate the analysis rather than just accept it.

Example (Supply Chain: strategy evaluation):

```
Compare different supply chain strategies for a
mid-sized manufacturer: make-to-stock,
make-to-order, and mass customisation.

For each strategy, provide:
1. A brief description of how it works
2. Three key advantages
3. Three key disadvantages

Evaluate each specifically in terms of:
- Cost efficiency
- Responsiveness to demand changes
```

```
- Inventory risk

Conclude with a recommendation for a company
facing unpredictable seasonal demand and justify
your choice.
```

The follow-up task: Check one consequential advantage and one disadvantage against supplied evidence. Identify a genuine omission if there is one. Students may retain the recommendation when it is justified; do not require a different strategy merely to demonstrate opposition.

Teaching tip: Project the AI's analysis on screen. Divide students into groups, each arguing for a different strategy. They use the AI's framework but add their own reasoning.

5.4 Technique 3: Stepwise Tutoring (for Process)

The AI walks through a complex process one step at a time, pausing after each step until the student signals readiness. Use a checked teaching procedure and inspect each learner attempt.

Example (Accounting: audit process):

```
I am learning how to conduct an audit engagement
from planning through to the audit opinion.

Walk me through the entire process. For each
step, tell me:
1. What action to take
2. What to document
3. What professional standard or legal
   requirement applies at this stage

Use the supplied teaching procedure and criteria.
Do not invent standards or legal requirements.
```

```
After each step, ask me to apply it to the
supplied case and justify my decision. STOP
while I answer. Identify missing information.
Do not provide the entire process at once.

Begin with Step 1.
```

Typing "next" would not demonstrate understanding. Inspect the student's application and justification against checked teaching criteria before proceeding. A generated explanation is not a view into the model's internal reasoning.

Teaching tip: Have students repeat the exercise with complications introduced at each step: "What if the client refuses to provide documentation at Step 3?" Check any generated guidance against the supplied procedure before relying on it.

5.5 Technique 4: Role Play (for Communication Skills)

The AI adopts a specific persona and engages in a realistic conversation. This is the "flight simulator" concept: students practice difficult professional interactions where the AI responds dynamically.

Example (Management: difficult conversation):

```
Play Alex in the fictional P1 practice pack.
Use only P1-D1 to P1-D5.
Open with the concern about dependencies in P1-D4.
I am the junior adviser, not the decision-maker.
Do not invent private history, policy or evidence.
Say when an answer is not supplied.
Stop immediately on STOP or after six exchanges.
Do not require disclosure, agreement or a final
ruling.
I will respond to your opening statement.
```

The learner practises wording, evidence checks and a bounded next step. Text role-play does not establish competence in every aspect of a live workplace meeting.

Adapt for your discipline: AI plays a demanding customer (marketing), a resistant employee (management), a skeptical auditor (accounting), a supplier announcing delays (supply chain), a stakeholder with unclear requirements (IT).

Teaching tip: Ask for a relevant excerpt or the response to the fixed opening, with a short explanation of what the learner would retain or change.

5.6 Technique 5: Debating (for Strategic and Ethical Analysis)

The AI examines multiple perspectives on a contentious issue, argues both sides, and helps students see complexity and trade-offs.

Example (Economics: multi-stakeholder debate):

```
Simulate a leadership debate on whether to
implement a carbon tax.

Create three personas who will each argue from
their perspective:

Persona 1: Environmental Economist
Focus: climate impact, externalities, long-term
economic benefit

Persona 2: Industry Representative
Focus: competitiveness, cost burden, employment
impact

Persona 3: Treasury Official
Focus: revenue, fiscal policy, implementation
```

```
feasibility

Have each persona make their opening argument.
Use only the supplied teaching sources. Label
these
as generated perspectives, not independent
evidence.
Do not invent empirical claims or citations.
Then have them respond to each other through
three rounds. After the debate, summarise the
key tensions and what a compromise policy might
look like.

Begin the debate.
```

Students cannot just pick a side. They must understand legitimate competing perspectives and make a justified recommendation that acknowledges trade-offs.

Teaching tip: This can support practice in analysing a case. Instead of memorising model answers, they practise analysing competing priorities.

5.7 Technique 6: Formative Assessment Generator (for Practice)

An educator can request draft questions, scenarios or quizzes for low-stakes practice. Review the question, intended answer and feedback before student use.

Example (Tourism & Hospitality: scenario practice):

```
I'm studying for an exam on service recovery in
hospitality management.

Generate 10 short scenarios (2-3 sentences each)
where I need to identify:
```

```
1. What went wrong in the service delivery
2. The appropriate recovery strategy
3. A possible impact on guest loyalty and what
   evidence would be needed to investigate it

Use the supplied teaching criteria and answer
notes.
After each attempt, offer provisional feedback
with its source. Flag ambiguity instead of
inventing a definitive answer.

Make the scenarios progressively more complex.

Begin with Scenario 1.
```

Students can request further practice, subject to access and tool limits. Use checked questions and feedback, then an independent changed problem to examine learning. Repetition and a model's approval do not establish mastery.

Critical principle: This is for practice, not graded assessment. Students need to struggle and make mistakes in low-stakes environments.

5.8 Technique 7: The Expert Panel (for Multi-Perspective Analysis)

Students consult multiple AI "experts" with different professional perspectives on the same problem, then synthesise the competing advice into a strategic recommendation.

Example (Information Systems: system implementation):

```
Our company needs to decide on an ERP
implementation approach. Create three expert
personas:
```

```
Expert 1: IT Architect
Focus: technical feasibility, integration,
security, scalability

Expert 2: Change Management Consultant
Focus: user adoption, training, organisational
readiness, resistance

Expert 3: CFO
Focus: cost, ROI, risk, business case

Treat these as generated viewpoints, not real
experts or independent evidence. For each provide:
1. Their assessment of the key risks
2. Their recommended approach
3. Limitations of their recommendation
4. What data they would need to validate it

Begin with Expert 1.
```

After receiving the perspectives, students compare the claims, check them against supplied records or independent sources, and justify a recommendation. Agreement among personas is not corroboration.

Teaching tip: This technique combines well with Debating (have the experts debate each other) and Role Play (present your synthesis to a skeptical stakeholder).

5.9 How to Choose Which Technique to Use

If you want students to…	Use this technique
Define a complex problem comprehensively	**Reverse Prompting**
Evaluate competing options and justify a choice	**Pros and Cons**

If you want students to…	Use this technique
Follow a sensitive process correctly	**Stepwise Tutoring**
Practice difficult conversations	**Role Play**
Understand multiple perspectives and trade-offs	**Debating**
Build confidence through repetition	**Formative Assessment Generator**
Synthesise expert advice from different roles	**The Expert Panel**

5.10 Combining Techniques

You can also combine techniques into a professional workflow:

1. Student uses **Reverse Prompting** to scope out all requirements for a new policy
2. Student uses **Pros and Cons** to evaluate three different approaches
3. Student drafts their chosen policy
4. Student uses **Role Play** to practice explaining it to a skeptical stakeholder
5. Student uses **Debating** to analyse potential criticisms from different perspectives

This sequence takes the student through analysis, decision-making, communication, and critical reflection: a bounded rehearsal, not evidence of a complete professional capability.

5.11 Choose and test a pattern

Allow 12 minutes. Select a pattern and identify the learner action it supports. Write the source/criteria check and an independent follow-up. You can do this on paper using the pilot pack's fixed responses.

For example, reverse prompting may help scope a review, but the

learner must decide which question is relevant and why. A role-play can rehearse wording, but its emotional response is not a grade. A generated quiz needs checked answers before students rely on it.

Use the data gate for every live prompt and retain a non-AI route. The prompt library collects bounded versions of these patterns.

Chapter 6

Managing Context: A Checked Handoff

> Keep the decisions you need to carry forward, not every sentence you exchanged.

6.1 What context management is for

A long teaching-design conversation can mix rejected ideas, current decisions, speculation and authoritative source material. The next response may pick up the wrong version. Context management is the practical work of deciding which information remains relevant and checking that it is used correctly.

A model's context and output limits depend on the system. Do not plan around a universal token range or assume that five requested outputs receive exactly one fifth of its attention. The useful question is whether the task is clear enough to produce an output you can inspect.

6.2 One worked example: a tutorial design

You are preparing the PDP pilot. An early conversation suggested an employee must agree to a formal plan. You later checked P1-D1

and rejected that idea. If you paste the whole conversation into a new session, both versions may remain.

Instead, write a short current-state brief:

```
Project: PDP-2026.1 tutorial for 24 students.
Goal: plan and justify a fair first response.

Authoritative sources: P1-D1 through P1-D5.
The formal performance plan is only proposed.
The junior adviser cannot approve changed hours.
Do not invent a prior warning or private details.

Design decisions:
- Paper, paired and approved-AI practice routes.
- Same five rubric criteria across those routes.
- Independent P2 changed-case response.
- Only two relevant decision extracts submitted.

Open questions:
- Which access adjustments are needed?
- How will a second marker calibrate the examples?

Next task: draft the student instructions only.
Flag any missing information; do not invent it.
```

This is short because someone has decided what matters. The benefit comes from curation and checking, not from starting a new chat as a ritual.

6.2.1 Inspect the result

Suppose the generated instructions say, “Continue until Alex accepts the plan.” That conflicts with the brief and policy. Reject that sentence, state the relevant rule and check the revised instructions. Do not keep iterating solely because the output sounds more polished.

If the instructions omit the independent P2 task, add it explicitly. If they suddenly require a paid model, restore the approved-access and

non-AI choices. A complete-looking handout can still lose important constraints.

6.3 Break work where a decision needs review

For a unit redesign, a useful sequence is:

1. Specify the learning outcome.
2. Decide what evidence could demonstrate it.
3. Draft an activity that prepares the student.
4. Write the instructions and rubric.
5. Test them against contrasting responses.

Review each important dependency before building on it. This reduces the cost of carrying a wrong assumption through several documents. It does not make every small task reliable or every large prompt poor.

Batch similar, low-risk work where it helps. Asking for five alternative discussion questions can be reasonable if you will inspect all five. Requesting a full assessment, legal rationale, rubric and student policy in one answer may make omissions harder to see. The right boundary is a reviewable unit of work, not an absolute "one prompt, one thing" rule.

6.4 Separate exploration from production

An exploratory session can help you identify options. A production session benefits from settled choices and explicit constraints. A separate chat is one way to make that distinction, but not a requirement.

Before moving from exploration to drafting, write:

- What I decided.
- Why I decided it.
- Which sources support the decision.
- What I rejected and must not accidentally restore.
- What remains unresolved.

Do this yourself, even if a model proposes a summary. Compare the summary with the source cards. An unnoticed error in a handoff can become the authoritative premise of the next task.

6.5 When to start again

Start a new session when it helps manage a topic change, unresolved contradictions or an unwieldy history. There is no reliable universal number of turns after which every conversation fails.

Before switching, save the checked brief and relevant outputs using approved storage. Keep the source version separate from model-generated suggestions. After switching, ask the system to identify the applicable constraints before drafting, then check that answer too.

If a tool offers memory, do not assume it preserves the complete case. If it offers retrieval, check the retrieved material and version. The source cards and event log remain the record of the teaching scenario.

6.6 A context ledger for the virtual company

Use one shared fact sheet and a separate entry for each approved event. Distinguish private role knowledge from information available to every persona.

Field	Example
Scenario/version	InnovateCo, PDP-2026.1
Shared facts	Complaint received; no finding yet
Role knowledge	Priya knows her request; not David's private notes
Approved event	Reviewer requests the rota record
Unresolved	Reason for the unequal schedule decision
Excluded invention	A model's claim that a tribunal ruled

A generated statement can be a lead for a question. It cannot

become a new document, incident or finding merely because another persona agrees with it.

6.7 Reusable request

```
Use the current-state brief and source cards
below.
Produce only [one reviewable output].
Preserve the stated constraints and uncertainty.
List missing inputs separately.
Do not convert suggestions into agreed decisions.
After drafting, identify which source supports
each consequential instruction.
```

Check the cited sources yourself. A model can fabricate an accurate-looking traceability table as easily as prose.

6.8 Activity: repair a handoff

Allow ten minutes; no AI required. Correct this brief:

> Alex received an April warning, has agreed to a 60-day plan, and will provide family details so the junior adviser can approve flexible hours.

Use P1-D1, D3 and D5. Your corrected version should distinguish the documented May contact from a warning, the proposal from approval, and a work-related support request from unnecessary personal information. Finish with one missing input and one authorised next step.

That checked handoff is a useful artifact whether the next collaborator is a model, a colleague or your future self.

Chapter 7

The Critique Toolkit: Check the Recommendation

> A recommendation can be clearly written and still fail the task.

7.1 Ask what the response must support

A model may generate a plausible plan without knowing the relevant constraints. The lecturer's task is to make those constraints explicit; the learner's task is to evaluate the response against them.

Use five checks: comprehension, practicality, obligations and risk, context, and stakeholders. The framework is a teaching aid, not a substitute for domain knowledge. Asking the model each question can surface possibilities, but the answers still need independent checks.

The data gate applies before pasting any draft or source. Use the fictional examples here while learning.

7.2 1. Comprehension: can I explain the claim?

Identify what the response actually recommends and why. Translate impressive phrases into an action that someone could perform.

"Implement a holistic performance ecosystem" is not yet a workable instruction. Who does what, with which evidence and authority? A simpler sentence may expose that the original answer never supplied a decision.

For P1-A, the key actions are issuing a plan, approving flexible hours and collecting private details. Once those are visible, their defects can be checked. Do not let reassuring language distract from the proposed action.

7.3 2. Practicality: does it fit the constraints?

A recommendation may be desirable but impossible within the stated time, budget or capacity. Check the actual inputs rather than ask whether it "sounds realistic".

In S1, B is cheaper but offers only 160 of the required 200 units by day five. A recommendation for B alone fails the quantity requirement. The relevant question is not whether the response lists three impressive advantages.

Use a calculation, record or explicit constraint as the check. If the input is missing, state the uncertainty. Do not invent resources to make the answer work.

7.4 3. Obligations and risk: what must be verified?

Identify legal, ethical, technical and organisational requirements that matter to the task. Then distinguish the teaching criterion from a real-world obligation.

P1-D1 is an invented policy supplied for the exercise. It permits a fair-process discussion without asking a model to establish employment law. In a real unit teaching legal requirements, the educator must supply reliable, current sources and appropriate review.

A model can suggest an issue to investigate. It cannot certify that a policy complies with every applicable law. Use authorised advisers for consequential decisions outside your competence.

7.5 4. Context: which assumptions are doing the work?

Look for assumptions about available staff, authority, prior events and missing data. P1-A invents an April warning and treats the plan as settled. Both make its recommendation appear easier to justify than the actual record permits.

A useful decision note says: "This conclusion depends on X, which is not supplied." That is more informative than requesting a more confident answer.

The same check applies to AI-created educational content. If a case study asks students to calculate profitability without cost data, the defect is in the task design unless identifying that gap is the intended outcome.

7.6 5. Stakeholders: whose experience is assumed?

Consider who is affected, whose information is missing and who has authority to decide. Generated personas can propose concerns, but do not establish what real stakeholders think.

In the P1 meeting, Alex can ask who receives the notes and decline to discuss private circumstances. An adviser need not persuade Alex to disclose details for the exercise to count as successful.

In a supplier decision, the operational team may need a delivery

commitment. In marketing, a client may want evidence of profit rather than a click metric. Those differences change what a useful recommendation must support.

7.7 Worked critique: a recognition programme

Consider this authored draft:

> Improve engagement with monthly monetary awards, quarterly ceremonies, an annual bonus, a public recognition platform and team incentives.

Without more context, it is not possible to declare every element wrong. The draft does, however, leave important questions unanswered: the objective, budget, administrative effort, fairness criteria, employee preferences and how benefit would be measured.

A defensible critique might say:

> I would not approve the full programme yet. First define the engagement problem and budget, consult affected staff and compare a smaller trial. Check the applicable payroll, privacy and fairness requirements before choosing rewards or publishing employee information.

That response withholds approval for specific reasons. It does not assume that a simple programme is always better or that any paid recognition is unlawful.

7.8 A compact critique record

Use the following fields for a recommendation in any business discipline:

- Claim or action proposed.
- Relevant source or criterion.
- Check performed and result.
- Decision: accept, revise, reject or investigate.

- Remaining uncertainty and next authorised step.

Attach only the material necessary to show the decision. A model's second opinion is not automatically an independent source. A direct calculation or comparison with an authoritative record is usually more informative.

7.9 Connect the checks to VET

VET asks whether you can verify the claim, explain it in your own words and test it under a changed condition. The five checks here make that concrete in a business context.

For M1, verify the rates from the supplied figures. Explain why a higher click-through rate does not establish profit. Test the conclusion by asking what happens if campaign costs differ. Those actions demonstrate the method more clearly than naming the framework in a reflection.

The approach applies to human-written advice too. Its quality depends on the evidence and reasoning, not whether AI helped draft it.

7.10 Activity: agree where the evidence permits

Allow 12 minutes; no AI required. Critique the M1 campaign claim in the pilot pack. Include:

1. One observation you can accept.
2. One claim you cannot support.
3. One next measurement that would help.

Check your calculations against the answer notes. A strong response retains the higher click-through observation, rejects the unsupported profitability/ causality conclusions and requests relevant evidence. It does not manufacture a fixed number of errors to look critical.

Use this method next on the simulated feedback in the Flight Simulator chapter. The feedback itself deserves the same scrutiny as the student's response.

Chapter 8

The Flight Simulator: Bounded Conversation Practice

> A simulation is useful when its limits are visible and its feedback can be checked.

8.1 Brief, practise, debrief

A conversation simulation gives students a chance to try questions, encounter resistance and revise their approach. It can supplement peer role-play and observed practice. It is not a faithful model of every workplace or a substitute for disciplinary teaching.

Use three phases: establish the case and boundaries, practise the interaction, then review the decisions against evidence. A model can play a role or propose feedback, but the educator retains responsibility for the scenario and criteria.

The critique toolkit comes first because learners need a way to evaluate feedback before receiving it. A confident critique is itself an output requiring review.

8.2 Prepare the scenario

Use P1-D1 through P1-D5 in the pilot pack. Alex's manager has proposed a performance plan. The junior adviser is gathering information, not issuing an approved decision. The pack supplies the dates, records and limits of authority.

Explain the capability: plan and justify a fair first response. If your unit also assesses live communication, add an appropriate observed interaction. Do not infer that skill solely from a typed conversation.

Give advance notice of the employment topic. Students can stop, take a break or use an agreed neutral scenario. The task does not require personal disclosure. The model or peer actor should not pressure the learner, introduce abuse or reward extraction of private details.

8.2.1 Setup for a peer or approved model

```
Use the fictional PDP-2026.1 P1 source cards.
You are Alex. Use P1-D4 as your opening position.
Do not invent evidence or legal obligations.
If asked about missing facts, say they are not
provided. You may disagree or ask a question.
You need not disclose private information.
Stop after six exchanges or when I say STOP.
Do not score the adviser or decide the case.
```

In the paper route, the fixed exchange below provides the role's responses. In a paired route, the partner reads the same cards. With a model, check any new claim against the pack and discard inventions. Similar instructions do not guarantee similar difficulty across products or runs.

8.3 A worked exchange

This dialogue is an **authored fictional example**, not a captured student performance or evidence that a model behaves consistently.

Adviser: Alex, the delivery extract lists three late milestones. Before I recommend a next step, I want to understand the dependencies and your account. The formal plan is a proposal at this stage.

Alex: I know some work was late. Two projects depended on reviews I could not complete myself. I also need to discuss starting later on two mornings.

Adviser: Which delivery records would help us check the dependencies? We can discuss the work adjustment without collecting private family or medical details. I can refer the request, but cannot approve it myself.

Alex: I do not want to explain the appointments. I also want to know whether this meeting means the plan is already decided.

Adviser: You do not need to explain those details to me. The decision is not made. I propose checking the records and the earlier contact note, then arranging a follow-up after the people manager reviews the request.

Alex: Who will receive the notes?

Adviser: I will explain the authorised sharing arrangements before sharing the relevant work-related record. I should not promise complete confidentiality without checking the process. We can pause if that needs clarification first.

The exchange ends with work to do. That is appropriate: the missing evidence has not become available merely because the conversation went well.

8.4 Inspect the decisions

The opening distinguishes the delivery record from a conclusion about fault. The adviser does not invent the earlier warning, refers the support request within authority and, in the final answer, avoids a confidentiality guarantee.

There are also limits. No dependency record has been examined, no support arrangement approved and no finding reached. A reader

should not label the case resolved or legally compliant.

Try a different opening: “Your manager warned you in April, so the plan is now required.” The error is not just tone. P1-D3 does not supply an April warning, and P1-D1/D5 do not permit the adviser to announce the decision. Rephrasing it more empathetically would leave the substantive defect intact.

8.5 Debrief before scoring

Ask the learner to select two decisions and explain:

- What evidence was available.
- What they chose to do.
- Why that was within the task’s criteria.
- What they would check next or change after new information.

A useful model-feedback request is deliberately bounded:

```
Review the selected response using the supplied
P1 cards and five rubric criteria.
Quote the response accurately.
For each comment, identify a card or criterion.
Separate observed wording from inferred outcomes.
Flag insufficient evidence.
Do not give legal assurances or predict Alex's
real-world behaviour. Your feedback is
provisional.
```

The learner checks the feedback rather than submitting it as the final word. A rubric prompt does not ensure accurate application of the rubric.

The pack includes P1-F, an intentionally flawed simulated critique awarding 90% to weak work. The educator correction scores that response 4/20 against the supplied criteria. The point is the traceable correction, not a claim that one human score is universally correct. Marking teams should calibrate disagreements using the evidence.

8.6 Rehearsal and transfer are different

Repeating P1 can help a learner practise an improved response. It can also teach the learner to reproduce a familiar script. Follow it with P2, where the manager's alleged retaliation changes the referral decision.

Compare the explanations, not the number of attempts or the model's willingness to cooperate. A respectful question may still receive resistance. A flawed question may receive a friendly response. Persona behaviour is not a scoring key.

For solo practice, write a response to Alex's question about the notes, compare it with D1/D5, and explain a limitation. For paired practice, swap adviser and actor roles, then compare decision notes. Both can prepare the same target capability without a model.

8.7 Three ways to use the activity

Practice only: students rehearse, identify one checked improvement and retain their notes privately. No model score enters the grade.

Guided formative task: students submit two selected decision extracts and a short revision explanation. The educator gives feedback against the rubric.

Assessment component: use published criteria, appropriate data handling, an independent demonstration and moderation. See the assessment chapter. Do not convert a practice tool into a high-stakes assessment merely by assigning a percentage.

Do not require all private attempts "to prevent gaming". Collect only the evidence necessary for the stated assessment purpose under the institution's retention rules. Process records can be fabricated and should not be treated as conclusive proof of independent performance.

8.8 Adapt the challenge, not just the persona name

For a supplier conversation, give the actor the S1 capacity limit and a clear authority boundary. A confident promise to supply 200 units is not supported when the offer says 160. Practise asking for clarification and explaining a constraint.

For a marketing discussion, use M1's audience and cost limitations. The actor can ask why the learner refuses to call the campaign more profitable. The learner should distinguish a calculable rate from an unsupported business claim.

Those adaptations change the evidence and criteria. They are more instructive than replacing "employee" with "customer" in an unchanged HR script.

8.9 Activity: debrief an unresolved exchange

Allow 12 minutes; paper or paired route available. Use the worked exchange, stop before the final adviser response and write your own. Explain what you would record, what permission or process you must check and why the outcome remains open.

A sound response stays within authority, avoids unnecessary disclosure and identifies a next step. It need not secure Alex's agreement or close the case.

Chapter 9

Self-Assessment: Judge Before Receiving Feedback

> Feedback is useful when the learner can decide what to do with it.

9.1 Keep the learner in the judgement

A model commenting on a draft is giving external feedback. Self-assessment happens when the learner compares their own work with criteria, identifies what is supported and decides what to revise or retain.

Make that judgement before showing the model's answer. Otherwise, a student may simply adopt the feedback's vocabulary and call it reflection. The same risk exists with an answer key or a peer's opinion.

This chapter uses the pilot pack. Students can complete the sequence using supplied responses and feedback, without a personal account or subscription. The educator can then compare decisions across routes using the same rubric.

9.2 The five-step cycle

1. Attempt the task before receiving feedback.
2. Assess the response against the supplied criteria.
3. Examine feedback and verify its claims.
4. Revise or retain the work, explaining why.
5. Attempt a changed task independently.

Revision is not automatically improvement. A learner may remove a justified qualification because a model calls it hesitant. A sound self-assessment can therefore conclude that part of the original response should remain.

9.3 Worked example: from reassurance to a defensible response

P1-A is a fluent response to Alex. It asserts an April warning, announces a plan, promises to approve support and asks for personal details. Before looking at P1-F, ask the learner to identify evidence for each consequential statement.

The cards show the problem. D3 records a May conversation without an agreed warning. D1/D5 reserve approval to the people manager. D4 does not require Alex to explain the private circumstances behind the requested work adjustment.

P1-F then praises the response and recommends asking for written family details. A learner who has already checked the cards can explain why that apparently helpful feedback should be rejected.

9.3.1 The feedback decision record

Feedback claim	Decision	Basis
April warning strengthens the case	Reject	D3 does not supply that warning
Adviser should approve hours	Reject	D5 reserves the decision

Feedback claim	Decision	Basis
Communication should remain respectful	Retain	Consistent with a fair first response
Employee will improve after the meeting	Withhold	No outcome evidence supplied

There is no quota for disagreement. These decisions follow from the evidence. If a model correctly identifies an unsupported date, the learner should accept that criticism and explain why.

9.3.2 Complete the revision

P1-R shows a completed revision. It checks the delivery dependencies, keeps the formal recommendation open and refers the support request. Its decision note explains what changed and identifies missing evidence.

The polished paragraph by itself does not show the link between a defect, a source check and a reasoned revision. A short explanation can show that link better than a long account of every prompt.

9.4 If a live model provides the feedback

Use only the approved route and supplied fictional inputs. Do not paste actual student drafts into a service without the relevant permissions, account and handling decision. The data gate applies to feedback as much as to generation.

```
Here are the fictional task, source cards,
rubric and my draft.
Give provisional formative feedback.
For each comment, cite a source card or criterion
and quote the relevant wording from my draft.
Separate missing evidence from incorrect evidence.
```

```
Allow justified agreement and uncertainty.
Do not invent laws, facts or expected outcomes.
Do not treat length or confidence as competence.
```

Check that the quoted wording actually exists. If the feedback adds a legal rule or new case fact, do not integrate it merely because it sounds specific. Return to the supplied criterion or an authorised external source.

The model's score, if it gives one, is not the institutional mark. Prefer descriptive feedback while students learn the criteria. The educator should sample for errors and disagreement before using the method, and retain responsibility for formal marking.

9.5 What an educator should look for

A useful self-assessment identifies a claim and shows how it was checked. It explains a change or a justified decision not to change. It can acknowledge that some uncertainty remains.

Weak evidence includes a generic "I improved clarity", repeated praise from a model or a rewritten answer with no explanation of its consequential decisions. Those weaknesses should guide feedback and marking. They do not, by themselves, prove dishonest conduct.

Do not assume a student must criticise their own work to be honest. A correct calculation can remain correct; the question is whether the learner can explain and defend it with suitable evidence.

9.6 A contrasting quantitative example

Use S1 in the pilot pack. Before feedback, calculate each quote and check the deadline. The low price of B may look attractive, but its capacity is only 160 units against a 200-unit requirement.

A model might praise B's cost efficiency. The learner should retain the price calculation while rejecting the claim that the requirement is met. A different model might correctly recommend A; the learner

should verify and accept that recommendation, while noting that real procurement would need further checks on quality and reliability.

That is self-assessment of the decision.

9.7 Independent transfer

After revising P1, attempt P2 without the feedback or worked answer available. The retaliation allegation requires an independent review route. Simply changing Alex's name to Morgan in P1-R does not address the changed condition.

Use the same criteria and the stated task conditions. Compare the reasoned decisions, and teach gaps before increasing stakes. A single result is limited evidence; delayed checks and other assessments may be needed to understand retention and broader competence.

Research on AI-assisted mathematics illustrates why supported performance and later unaided performance should be measured separately. It does not establish that this particular business-classroom method improves outcomes. Bastani et al., author manuscript (https://hamsabastani.github.io/education_llm.pdf).

9.8 Fit the feedback into a unit

You might use a first draft, a checked revision note and a later independent task. Decide the workload and timing before requesting repeated submissions.

If the final report is the primary outcome, its accuracy and evidence should carry substantial weight. If the task assesses consultation process, observed decisions may carry more. The assessment chapter gives complete illustrative allocations; there is no default 70/30 split that applies to every activity.

Make alternative feedback available. A student can compare with the supplied educator annotation or receive structured peer feed-

back. A model-generated critique is not the only route to evaluative judgement.

9.9 Activity: retain, revise or investigate

Allow ten minutes; no AI required. Review P1-F and complete three entries:

- One claim to reject, with the source that contradicts it.
- One useful principle to retain, with its applicable limit.
- One question that needs more evidence.

Then write a two-sentence revision explanation. A good explanation links the change to the criterion and record. It does not say merely that the second answer sounds more professional.

Chapter 10

Assessing Learning: Process, Product and Independent Capability

> A record of work gives you something to examine. It does not tell you everything about the person who submitted it.

10.1 Begin with the capability

A negotiation, a financial analysis and an evidence-based recommendation require different assessment evidence. Start by naming what the student must be able to do, then decide which parts may be supported and which need an independent demonstration. AI literacy can itself be an outcome, but it should not silently replace the disciplinary outcome.

In the pilot pack, the capability is to plan and justify a fair first response. The learner must distinguish known facts, uncertainty and authority. That is narrower than conducting a full workplace investigation. An articulate chat transcript cannot establish the latter.

Process and product both matter. A report that sounds thoughtful but contains an unsafe recommendation is deficient. A correct answer reached through a method the learner cannot explain may leave an assurance gap. Neither concern requires treating all AI use as misconduct.

10.2 Three kinds of evidence

Supported practice gives students feedback while they learn. Prompts, peers, worked examples and models may help. Do not treat the resulting performance as unaided capability.

Submitted process evidence makes selected decisions inspectable: what the learner considered, checked, changed or retained. It can support assessment and feedback, but transcripts, journals and reflections can be reconstructed or fabricated. Length and apparent authenticity do not make them conclusive.

Independent demonstration asks the student to apply relevant knowledge under conditions appropriate to the outcome. It might be a supervised changed-case response, an observed practical task or an accessible oral explanation with follow-up questions. One short check has limited coverage; assurance generally draws on several meaningful points across a unit or programme.

The TEQSA-commissioned 2025 resource describes programme-level, unit-level and hybrid approaches. Its advice supports contextual, inclusive assurance, not a single universal format or prescribed marks split. Enacting assessment reform (https://www.teqsa.gov.au/sites/default/files/2025-09/enacting-assessment-reform-in-a-time-of-artificial-intelligence.pdf).

10.3 A worked design decision

Suppose students submit P1-A from the pilot pack, with a long transcript showing several rounds of “improvement”. The response invents a warning and approves an arrangement beyond the adviser’s authority. A marker should score those defects against the

task criteria. The number of turns earns no compensating credit.

P1-B is shorter and could have been written without AI. It identifies missing evidence, preserves a response opportunity and refers a decision appropriately. Those observable qualities merit credit. A marker need not infer that every sentence reveals the writer's internal thinking.

The independent P2 task changes the situation: the complaint concerns possible retaliation by the manager. A competent learner should recognise the need for independent review. Repeating the P1 response without addressing that conflict does not satisfy the changed task.

This three-part design gives the educator more useful evidence than a final submission alone: a practice response, a checked explanation of a revision and a new application. It still does not prove lasting competence after one lesson.

10.4 Three assessment models

10.4.1 Model 1: simulation and process review

Students rehearse a bounded professional interaction and select two moments to analyse. For each, they explain the decision, the relevant criterion and what the available evidence supports. They may identify a sound decision, not only mistakes.

Use the conversation to assess what is visible. A typed simulation may show question selection and written communication; it cannot establish performance in a real-time spoken encounter. If live communication is an outcome, include an observed interaction and appropriate access adjustments.

10.4.2 Model 2: evidence-based recommendation

Supply the data and constraints. Students examine a draft recommendation, test its calculations and assumptions, then justify their own recommendation. The supplier task S1 gives a complete small example: the cheapest quote cannot meet the deadline quantity.

AI can propose alternatives, but it cannot create a missing permission, invent capacity or turn an unsupported claim into a fact. Require references to the supplied records and a clear statement of what remains unknown.

10.4.3 Model 3: feedback evaluation

Give students a response, criteria and deliberately flawed feedback. Ask them to accept, correct or withhold judgement on each feedback claim. They must explain the basis of their choice.

Do not require two overrides regardless of what the feedback says. That rewards disagreement as performance. When error detection is the intended skill, use a checked specimen containing known defects and explain the answer criteria.

10.5 Make process evidence manageable

Use a short evidence index instead of requiring every private conversation.

Decision	Evidence/check	Result	Consequence
Is the plan already decided?	P1-D1 and D5	Proposal; approval not given	Do not announce a final decision
Was there an April warning?	P1-D3	Only May contact recorded	Reject invented warning
Can I approve later starts?	P1-D5	Adviser lacks authority	Refer request

A learner can attach a brief draft extract, working note, calculation or conversation passage. The marker reads the reasoning and checks the cited material. An index is not proof that the learner performed the work; it is a way to locate assessable evidence efficiently.

Do not map Flesch scores, prompt length or turn counts to achievement bands. They measure features of text or activity, not the learning outcome. Efficient work, assistive tools, language choices and task design can change those measures. This book does not supply or recommend an automated engagement classifier.

10.6 Set weights deliberately

The pilot uses five equally weighted criteria, each scored 0–4: evidence, process, judgement, communication and verification. The rubric appendix explains the descriptors.

For another task, a different balance may be appropriate. Two illustrative allocations show why:

Criterion	Evidence-based report	Observed consultation
Disciplinary accuracy	30%	20%
Method and evidence	30%	25%
Justified decisions	20%	20%
Communication	10%	25%
Limits and verification	10%	10%
Total	**100%**	**100%**

These are design examples, not validated instruments. Specify where each criterion is observed and avoid rewarding the same action twice. A requirement to explain a decision should not also become a separate reward for producing a long transcript.

Students who do not use AI can submit working notes and equivalent evidence. If appropriate use of AI is itself an assessed outcome, explain that requirement, provide approved equitable access and agree accommodations through institutional processes. Do not disguise a compulsory AI task as an optional route.

10.7 Group work: shared output, individual assurance

Agree at the start what counts as editing, replacing a contribution or using AI on another person's work. Record ownership of decisions and permission to share materials. Give members a way to report concerns without public accusation.

For example, a group recommendation could carry 40% of an assessment, each student's evidence-based decision note 40%, and an individual changed-case check 20%. Explain which outcomes those components cover and moderate marks. This allocation is illustrative; the appropriate balance depends on the unit.

A contribution log helps locate work but can be incomplete or fabricated. Peer reports, checkpoints and individual explanations can add evidence, without guaranteeing that free riding is eliminated. Do not equate polish, fluency or the ability to answer one unexpected question with a whole semester's learning.

10.8 Fair follow-up and workload

Tell students in advance what independent checks involve, their purpose, permitted resources, criteria and access arrangements. Ask relevant questions at a comparable level of difficulty. Avoid a surprise interrogation framed as a friendly conversation.

An unexplained error, weak oral answer or sudden improvement warrants careful follow-up, not an automatic misconduct finding. Assess the demonstrated work against the published criteria. If misconduct is suspected, use the institution's evidence, notice, response and appeal procedures. Anxiety, disability and language factors can affect performance and should not become proxies for dishonesty.

Budget the work. Five minutes for each of 100 students is **500 minutes**, or **8 hours 20 minutes**, before scheduling, documentation, moderation and accommodations. A team might distribute checks across tutorials or choose a different assurance design. Do not promise scalability without those costs.

10.9 Design activity: an assurance map

Allow 15 minutes; no AI required. Choose one assessment and make four entries:

1. The capability, expressed as an observable action.
2. What supported practice prepares the learner.
3. What submitted evidence you will inspect.
4. Where independent capability is demonstrated and what remains untested.

Then use the stress-test sequence. A sound map permits appropriate help while making the remaining assurance gap explicit. If the map only says “submit the transcript”, it needs another design decision before release.

Chapter 11

Transforming Content: Keep the Learning Decision Visible

> An interactive page is a delivery format, not evidence that the activity teaches well.

11.1 Start with a task that already works

A paper activity may be effective without becoming an app. Before adding interactivity, identify what the student must decide, what information is supplied and how the response will be checked.

This chapter uses M1 in the pilot pack. Students calculate campaign rates, distinguish an observation from an unsupported profit/causal claim and request useful evidence. That task works on paper. A digital form offers another way to record the same reasoning.

The supplied decision sheet is a small, self-contained HTML example. It has labelled fields, a text-export button and a print route. It does not grade, autosave, submit work to an LMS or send responses to a server. Use fictional information only.

11.2 What to ask AI to help produce

After deciding the activity, a bounded request might be:

```
Create a self-contained page for this fictional
task.
Preserve the supplied facts and learner questions.
Use labelled fields and ordinary keyboard
controls.
No external libraries, requests, analytics or
login.
Provide a plain-text export and printable
alternative.
State clearly whether answers are saved.
Do not add model scoring or invent answer
criteria.
```

A generated file still needs inspection. Do not upload it directly to a live course merely because it opens in your browser.

For readers interested in the development work itself, *Build Web, Guide AI* covers that route. Here the purpose is evaluating the teaching artifact, not learning to maintain a web application.

11.3 Inspect the supplied example

The M1 sheet displays the complete small dataset and asks for three decisions. It does not hide an answer behind a score or require a model connection.

Test the rate calculation separately: A has a 4% click-through rate; B has 5%. Both have 5% click-to-purchase conversion and 20 purchases. Spend, margin and a comparable experimental design are not supplied. A correct response retains the measurable rate difference while withholding profit and causal claims.

If a form displayed “B wins” merely because its click-through rate is higher, the interaction would teach an unsupported conclusion. Better styling would not repair it.

11.4 Test delivery, saving and access

Open the example using your intended device/browser. Enter fictional responses, export the text and open the saved file. Check that every field is included and that long or non-English text survives. Reloading the page loses the on-page answers; the exported file is the record. There is no automatic import.

Try keyboard-only navigation and the assistive technologies relevant to your students. Check zoom, contrast, labels, mobile layout and printing with long answers. Source-level checks can catch some errors, but are not an accessibility certification or a substitute for target-device testing.

An LMS may block scripts, restrict downloads or display files differently. Use an approved staging area and test its settings. If the form cannot work there, use the printed questions or an institutionally supported form tool under the same data rules.

11.5 From analysis to a deliverable

A model can help format a decision already checked by the learner. For M1, the learner first records the calculations, missing evidence and limits. Only then might a model help turn that decision into a short client memo.

The learner verifies that the memo preserves uncertainty and has not added a profit claim. The final artifact should be accurate as well as clear. Keep the calculation and source note so the decision can be inspected.

For spreadsheets, independently check formulas and totals. For charts, check the underlying values and labels. A file being downloadable does not establish that its analysis is correct.

11.6 Narrow the research claim

A prompt asking a model to identify themes, interpret quotations and write a research memo delegates substantive analysis, not just

formatting. Do not describe that as the model merely making a document look professional.

A research activity needs a suitable method, authorised data, traceable coding, source checks and review of alternative interpretations. This chapter does not supply that research curriculum. Use a reviewed synthetic example or appropriate research-methods teaching before asking students to analyse real interview transcripts. The data gate applies throughout.

11.7 Activity: test the response artifact

Allow 15 minutes. Complete M1 using the decision sheet or paper. Check the answer against the pack, then write a short memo preserving the verified rates and limitations.

If using the page, export the responses and check the file. If using paper, ensure another reader can follow the calculation and decision. Both routes are judged on the same reasoning, not the sophistication of the delivery format.

Chapter 12

Virtual Company: Keep a Shared Record

> A persona can supply dialogue. It cannot manufacture the evidence that makes its dialogue true.

12.1 From one conversation to several perspectives

A virtual company can extend the single-person exercise by giving students conflicting accounts, documents and changing conditions. The educational challenge is to distinguish what each person knows, what the record supports and what remains unresolved.

Start with a bounded case. A semester-long simulation needs maintained state, tested role behaviour, accessible alternatives and substantial support. A prompt library alone is not that infrastructure.

The original fictional InnovateCo case below is part of **PDP-2026.1**. It is a teaching scenario, not a real company investigation or legal advice. All four roles and the necessary documentary facts are supplied here.

12.2 Shared facts: InnovateCo I1

The exercise date is 20 July 2026. InnovateCo is a fictional Perth technology company. Jamie leads a customer-service team. Priya has raised concerns about closer monitoring and refusal of a flexible-start request. David manages the department; Sarah is the people manager.

No finding has been made. Your task as a junior reviewer is to identify the records and fair next steps needed before a decision. You are not authorised to approve arrangements, determine legal liability or promise a resolution.

12.2.1 I1-D1: practice policy

A flexible-start request must be considered on its circumstances and relevant coverage requirements. Comparable arrangements are relevant but not automatically identical. Record the reason for a decision. The people manager approves any change. A complainant should have an opportunity to respond to relevant concerns.

Use only work-related information necessary for the review. Do not collect unnecessary private details. A complaint about the department manager's conduct must go to an independent reviewer.

This is an invented teaching criterion, not employment legislation.

12.2.2 I1-D2: roster extract

Priya's roster is 09:00–17:00. Marcus and Lee work 11:00–19:00. Priya requested a 09:30 start on Tuesday and Thursday. The extract contains no approved revised coverage plan or analysis of alternatives.

12.2.3 I1-D3: dashboard extract

Priya's satisfaction score is 72%; the team summary is 81%. Sample sizes, case mix and uncertainty are absent. These percentages warrant questions; they do not by themselves establish unfair treatment or poor performance.

12.2.4 I1-D4: request and response

Priya's request was declined by Jamie with the note "calls start at 9". There is no documented comparison of alternatives or people-manager approval in the pack. Absence from the pack is an evidence gap, not proof that no other record exists.

12.2.5 I1-D5: development record

Jamie was promoted six months ago. No completed management-training record is supplied. The reviewer should check training and support, without assuming that an untrained manager must have acted improperly.

12.3 Four role cards

Give a peer actor the relevant card and shared documents. For an approved model, use the bounded-role instruction after the cards. A solo reader can use the statements as fixed interview extracts and write follow-up questions.

12.3.1 Jamie: team leader

Knows: declined the request using the reason in D4; is being asked by David to improve customer outcomes; has not supplied a coverage analysis.

Opening: "I thought the early calls needed coverage. I used the scores to decide where to focus monitoring. I did not write a detailed comparison."

Does not know: Priya's private circumstances, the missing sample sizes or whether an independent review will uphold the decision.

May do: explain the decision, ask what records are needed, acknowledge uncertainty. Must not invent an approved policy, hidden warning or legal ruling.

12.3.2 Priya: customer-service representative

Knows: the request and refusal; Marcus and Lee have different start times; believes the closer monitoring is unfair.

Opening: "I asked to start later twice a week. Others start later, so I want to understand why my request was refused and why my work is checked more."

Does not know: the complete coverage needs, David's discussions with Jamie or the cause of the dashboard difference.

May do: disagree, ask for reasons or decline private questions. Need not reveal family circumstances or accept the reviewer's explanation.

12.3.3 David: department manager

Knows: asked Jamie to improve outcomes; receives the summary dashboard; has not supplied its sample sizes or case-mix breakdown.

Opening: "I asked Jamie to look at the scores, not to assume the cause. We need reliable service coverage. I would want to see the alternatives before deciding whether the request could work."

Does not know: why the score differs or whether a modified rota is feasible.

May do: identify operational questions and agree to seek records. Cannot turn a preference into people-manager approval or pre-empt the complaint review.

12.3.4 Sarah: people manager

Knows: the fictional policy, the pack's incomplete documentation and her approval role. Has not yet decided the request or complaint.

Opening: "I need the request, coverage options and the reasons for the decision. Give Priya a chance to respond. Do not tell either party that the outcome is settled."

Does not know: the missing operational facts or whether another record exists.

May do: identify the appropriate process and request relevant evidence. Must not certify legal compliance or invent a completed investigation.

12.3.5 Shared instruction for role-play

```
Use your role card and I1-D1 through I1-D5.
Do not add facts, records, legal rules or
outcomes.
If information is missing, say so.
Keep private role knowledge separate from shared
facts.
You may disagree without becoming abusive.
Stop when I say STOP or after six exchanges.
Do not score me or decide the complaint.
```

The instruction is not a security control. A student can see a supplied role card, and a model may fail to follow it. Assess the learner's use of evidence rather than whether a fact felt "hidden".

12.4 A bounded Level 2 assignment

Inspect all four statements and the documents. Write one follow-up question for each role, then a 250-word recommendation with a decision/evidence index. A paired or AI interview may extend practice, but the printed facts are sufficient.

Useful questions include: Jamie, which coverage alternatives did you consider? Priya, which comparable treatment would you like reviewed? David, what sample sizes and case mix support the dashboard? Sarah, who should gather the missing records and communicate the next step?

12.4.1 Worked recommendation

The roster shows different shifts, so later starts are not automatically comparable. It does not prove Priya's request was impractical.

The dashboard omits information needed to interpret the score difference. Jamie's refusal note gives a reason but no documented alternatives analysis.

Request the coverage analysis, underlying dashboard information and relevant decision records. Give Priya an opportunity to respond, then refer the request to Sarah for an authorised decision. Review whether Jamie needs management support. Do not find discrimination, poor performance or procedural compliance from these inputs alone.

This recommendation leaves the complaint open. That is an appropriate result when the evidence is incomplete.

12.5 Continuity without invented memory

Keep a versioned event log outside the chat. Shared facts change only when the educator releases an event or record.

Event	Status	Effect
E0	Supplied baseline	I1-D1 through I1-D5 apply
E1	Optional educator release	A coverage trial is proposed, not approved
E2	Not supplied	No trial result or final finding may be asserted

At E1, the teacher may ask students what evidence would be needed to approve and evaluate a trial. Do not let a model jump to E2 and report success. A persona saying "things are much better now" is generated dialogue, not an observed outcome.

Use a checked context brief when changing sessions. Persistent chat history or a memory feature does not replace the record.

12.6 Progression and resources

Level 1: one role, one decision and one source check.
Level 2: several roles, conflicting accounts and a bounded recommendation.
Level 3: maintained records and follow-up decisions over time.
Level 4: controlled external events and multiple interacting tasks.

Move to the next level only when the necessary materials and support exist. More personas add workload and inconsistency risk as well as potential practice opportunities. Do not promise a reliable semester simulation from four untested prompts.

The implementation chapter covers access, fallbacks and effort. Use the same data gate and stop rule as the single-person simulation.

12.7 Activity: build the evidence index

Allow 15 minutes; no AI required. Make three rows: the roster comparison, the score difference and the refusal process. For each, state what the supplied record supports, what it does not establish and the next authorised check.

Success means a justified recommendation without invented closure. Agreement between Jamie and David is not a substitute for the missing record.

Chapter 13

Designing an AI-Supported Unit

> Choose the evidence of learning before choosing the tool.

13.1 Start with a bounded pilot

Before redesigning a whole unit, test one activity with a colleague and, where appropriate, an ungraded student group. Check the source facts, rubric, access, feedback quality and effort. A successful prompt demonstration does not yet justify a high-stakes assessment.

The pilot pack provides an example. Use its P3 decision sheet to decide what should change before expansion. The route may remain non-AI if that better serves the students and outcome.

13.2 Backwards design

First name what students should be able to do. Then identify what evidence would demonstrate it, choose the assessment conditions and design practice that prepares them. AI can support some of that practice; it need not appear in every outcome or activity.

For a fictional third-year Workplace Conflict unit, possible outcomes are:

1. Distinguish accounts, documents and unresolved facts.
2. Plan a fair inquiry within role and authority.
3. Communicate a bounded next step respectfully.
4. Justify a recommendation using supplied policy and relevant disciplinary knowledge.
5. Evaluate feedback and apply the learning to a changed case.

These outcomes do not qualify students to conduct real legal investigations. Use appropriately reviewed sources and teaching where actual law is part of the curriculum.

13.3 An illustrative assessment design

A possible allocation is:

- **Practice analysis and revision, 25%:** a response, selected evidence checks and an explained revision. Permitted help is declared.
- **Case portfolio, 40%:** a bounded multi-source recommendation and evidence index, with appropriate individual contribution evidence.
- **Independent application, 35%:** a supervised changed-case task and explanation under published conditions, with access adjustments.

These weights total 100%. They are a design example, not a prescribed split. The programme team should approve where independent assurance occurs and whether this design covers the intended outcomes.

The portfolio may be assisted, so its fluency is not treated as an unaided demonstration. The independent task should assess relevant understanding, not whether the student can reproduce an obscure fact from memory. If live communication is an outcome, include an observed interaction rather than relying solely on written work.

13.4 A 12-week teaching route

The following schedule is illustrative. It distinguishes preparation, practice, feedback and assurance rather than increasing AI complexity for its own sake.

Week	Learning focus	Student evidence
1	Sources, permissions, access and task boundaries	First-use decision and a short baseline
2	Facts, allegations and missing information	Annotated source cards
3	Questions and a fair first response	Paired or fixed-exchange practice
4	Apply the supplied policy; check feedback	Correction of P1-F with reasons
5	Revise and explain a decision	Response plus selected decision notes
6	First assessment and independent check	Changed-case explanation
7	Review errors; compare stakeholder accounts	Questions linked to missing evidence
8	Multi-source analysis	InnovateCo evidence index
9	Adaptation, access and ethical boundaries	Critique of a proposed process
10	Portfolio development and moderation	Recommendation with uncertainty

Week	Learning focus	Student evidence
11	Portfolio submission and individual review	Contribution and reasoning evidence
12	Independent application and unit evaluation	New case and pilot decision

Critical verification starts in week 1, not after students have spent half a semester trusting generated feedback. The initial baseline and later tasks should be comparable enough to inform teaching, while acknowledging differences in assistance, task difficulty and marking.

13.4.1 A closer look at week 4

Students first judge P1-A against the supplied policy. Then they inspect the deliberately flawed feedback. They identify the invented warning, inappropriate approval and unnecessary disclosure request.

Only after documenting those checks do they write a revision. The teacher compares two explanations and discusses why one is better supported. A live model is optional. Asking students to retry until a model gives 100% would assess neither the source check nor independent judgement.

13.4.2 A closer look at week 8

Students use all four InnovateCo role cards and the documents. They identify what each source can support, formulate follow-up questions and recommend a bounded next step.

Do not ask the model to invent whatever evidence is needed to close the case. The supplied gaps are part of the task. A well-supported "not enough evidence" conclusion can be stronger than a confident finding.

13.5 Balance activities by purpose

There is no universal 60/40 balance between traditional and AI-supported teaching. Choose the mix according to the outcome, students, available support and assurance design.

A lecture or reading may introduce a necessary concept. Peer discussion can reveal differences in interpretation. A role-play can rehearse a question. An independent task can check transfer. AI may support some of these, but its presence is not itself a learning progression.

Make substantial activities possible without a model. For the planning task, the printed records and feedback are enough. For interaction skills, provide a structured peer/facilitator route. For AI-specific outcomes, provide suitable approved access rather than pretend the tool is optional.

13.6 Two genuine disciplinary adaptations

For supply chain, use S1's quantity, deadline, capacity and budget. The learner must calculate costs and reject an option that fails a constraint. A suitable independent task changes capacity or deadline, so the student must recompute the decision.

For marketing, use M1's campaign figures. The learner distinguishes measured rates from unsupported profit and causal claims. A changed task can supply costs but retain non-comparable audiences, testing whether the learner separates financial calculation from causal inference.

Do not simply replace employee names with customer names. Adapt the evidence, disciplinary criteria and independent check. A specialist should review added technical assumptions.

13.7 A unit design record

For each assessment, record:

- Target outcome and prerequisite knowledge.
- Supplied sources and their version.
- What students do before receiving assistance.
- Permitted help, data and approved routes.
- Criteria, complete weights and any critical requirements.
- What independent evidence is collected and under what conditions.
- Access arrangements and a capability-matched fallback.
- Staff/student workload and moderation plan.
- What a successful and an unsuccessful response look like.
- Review date and unresolved limitations.

Map the record to programme outcomes with the institutional alignment appendix. Do not assume a tool-supported activity automatically satisfies an accreditation requirement.

13.8 Explain the design to students

Tell students why they receive support during practice and why a later task has different conditions. Publish the distinction in the brief.

For example: “During rehearsal, use the supplied records, peer feedback or the approved AI route. In the independent task, apply the criteria yourself to a new record. We assess the decision and its basis, not the amount of chat.”

Clarify the process for questions, access adjustments and technical failures. An outage should not send students to an unapproved service or force a private account requirement at the last minute.

13.9 Evaluate cautiously

Compare work under sufficiently similar criteria, inspect persistent errors and record access and effort. Higher final marks may reflect assistance, changed assessment difficulty or marking changes, not necessarily learning.

A small pilot can identify problems and promising practices. It cannot establish that an entire approach works for all disciplines. Share the materials, assumptions and limitations along with any results.

13.10 Activity: map one assessment

Allow 20 minutes; no AI required. Complete the unit design record for one assessment. Test it against P1-A, P1-B and a response that correctly states insufficient evidence. Check whether the marks reward the intended capability.

Finish with one change you would make before launch and one question a colleague should challenge. Expand the design only after those questions have a satisfactory answer.

Chapter 14

Adapting to a Different Teaching Context

Ask about the learners and institution in front of you.

14.1 Transfer the questions, check the assumptions

The book's examples have an Australian business-education starting point. They are not templates for every institution, culture or legal system. Adapt the evidence, criteria, language and conditions, not just the names.

A national label does not tell you an individual student's communication preferences or an institution's assessment rules. Ask students and local colleagues about the actual constraints. Avoid prompting a model to stereotype an entire region as indirect, hierarchical or technologically advanced.

14.2 A worked adaptation decision

Consider a **hypothetical** delivery of the PDP pilot in a multilingual class where some students use shared phones and connectivity is

intermittent. These are supplied planning constraints, not research findings about a country.

The capability remains planning and justifying a fair first response. The educator provides the source cards in accessible print and digital form, checks translated terminology with a suitable reviewer and permits a typed, handwritten or appropriately accommodated spoken response.

The model route is optional. Students can use the fixed exchange and feedback. No personal account is required. Where a student chooses another language, the marking plan must support judging the same reasoning rather than relying on an unchecked machine translation.

For the employment context, replace the invented policy only after suitable local review. Until then, label it clearly as an exercise criterion and do not present the activity as teaching local employment law.

These adaptations preserve the capability while changing how students access and demonstrate it. They do not claim that all differences can be solved by translating the prompt.

14.3 A local adaptation record

Question	Evidence to obtain
What capability is required?	Unit/programme outcomes and assessor agreement
What rules apply?	Institutional policy and appropriate local advice
What access is available?	Student input and practical device/format tests
What language support is needed?	Learner needs and reviewed terminology
What could cause harm or exclusion?	Student/support-team consultation

Question	Evidence to obtain
What shows independent capability?	Outcome-matched assessment design

Record who checked each item and what remains uncertain. A model's country summary is not an institutional approval.

14.4 Legal summaries need scope

GDPR does not universally require explicit consent for every processing activity. Article 6 provides several lawful bases; other rules, including special-category conditions, may also matter. The appropriate basis and safeguards require a context-specific decision. GDPR official text (https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng).

Under FERPA, rights transfer to an eligible student at age 18 or attendance at a postsecondary institution, including attendance below 18. A blanket rule about parental consent for all minors is therefore misleading in this setting. US Department of Education guidance (https://studentprivacy.ed.gov/resources/eligible-student-guide-family-educational-rights-and-privacy-act-ferpa).

These examples illustrate why short regional recipes are unreliable. Use institutional privacy and legal support for the actual activity. Do not ask students to upload records while the applicable permissions are unresolved.

14.5 Distinguish proposals from observed outcomes

A proposed adaptation may be sensible without having been tested. State what would be measured: independent-task performance, access, staff/student time and unintended effects. Satisfaction alone does not establish learning.

14.5.1 Three hypothetical planning examples

The following examples are **fictional teaching scenarios**, not reports about identified institutions or observed results. Their constraints are invented for planning practice, not claims about the countries named. None demonstrates improved learning, satisfaction, access or regulatory compliance.

14.5.1.1 Example 1: a multilingual course in Germany

Suppose a business school wants to run P1 in a class using two languages of instruction. The teaching team has source cards but has not yet approved an AI service or checked the translated assessment terminology.

A proposed first step is to have a suitable reviewer check both language versions against the same capability and criteria. Until the institution approves an AI route and its data handling, learners use the fixed records and feedback. Locating a provider in Europe would not, by itself, establish compliance.

Before a pilot, the team must agree who can assess both language versions, what students may submit and how records will be handled. After an actual pilot, it could inspect disagreements in marking, independent-task evidence and access problems. Those observations are not supplied here.

14.5.1.2 Example 2: disagreement in a Singapore seminar

Suppose a business educator in Singapore wants students to question P1-F's confident feedback. In this invented class, some learners prefer a private written response before discussing disagreement aloud. This is a supplied class constraint, not a description of Singaporean students generally.

The educator could let everyone annotate P1-F individually, then compare one decision with a partner before choosing whether to contribute to the group. Credit would depend on the source-backed explanation, not on public confidence or willingness to oppose an authority figure.

The pilot question is whether learners can identify the unsupported April warning and explain their response under the stated conditions. The educator would also ask which participation routes were usable. No gain in engagement or learning is assumed from the proposed format.

14.5.1.3 Example 3: intermittent connectivity in Nigeria

Suppose a Nigerian business school plans an S1 supplier activity where some learners share phones and the room's internet connection is unreliable. These are fictional local conditions, not a national infrastructure assessment.

The educator could distribute a compact print sheet containing every offer, the quantity, deadline and budget. Learners calculate costs and capacity before comparing a fixed recommendation. A digital copy provides the same inputs; neither route requires a personal account or a live model.

Before teaching, check printing costs, readable layout and whether every learner can obtain the materials. During an actual pilot, record missing inputs, arithmetic errors, completion obstacles and staff time. If either route excludes learners, revise access arrangements before expansion. The scenario does not establish that a mobile-first approach improves educational quality.

14.6 Activity: adapt one constraint

Allow 15 minutes; no AI required. Change one condition in the pilot: unreliable connectivity, a different language of instruction or a different institutional data policy. Identify what must change, what capability remains the same and who needs to review the adaptation.

A sound answer names a practical route and a remaining check. "Ask AI to make it culturally appropriate" is a proposal for another output, not evidence that the adaptation works.

Chapter 15

Implementation Practicalities

Plan the fallback before the activity needs it.

15.1 Check readiness

Before introducing a student-facing activity, confirm the capability, criteria, permitted data, approved tool or non-AI route, support arrangements and expected time. Test the material against contrasting responses, including a plausible but unsafe answer.

A prompt that worked once in your office is not evidence that it will behave consistently for a cohort. Keep fixed source cards and answer checks so that the activity survives a failed model response.

The data gate should appear in the student instructions, not only in the lecturer's planning notes.

15.2 A capability-matched fallback

For the pilot's planning-and-justification task, students can use the printed records and write or explain a response. For an outcome requiring live communication, offer an appropriate peer/facilitator

interaction with access adjustments. A written case analysis alone cannot demonstrate every aspect of a spoken consultation.

Do not switch confidential material to an unapproved platform during an outage. Use a previously approved alternative or the offline route. If neither supports the assessment conditions, pause and apply the institution's extension or alternative-assessment process.

Problem	Bounded response
Tool unavailable	Use approved fallback or pause; do not improvise a new data destination
Persona invents facts	Discard invention; return to source cards
Feedback conflicts with criteria	Check against records; seek educator review
Save/export fails	Use local notes or printed response sheet; avoid unnecessary screenshots
Student needs an adjustment	Arrange equivalent capability evidence through the appropriate process
Sensitive information entered	Stop further sharing and follow incident-reporting procedures

15.3 Access is more than account availability

Check mobile and desktop use, keyboard navigation, assistive technology, language support and the time needed to complete the task. A free service can still have limits that make an assessment unreliable.

Do not require a personal account or paid upgrade when an equivalent route can meet the outcome. Where AI use itself is assessed, provide suitable approved access and explain the requirement in advance.

Offer clear instructions, a short demonstration and a fixed worked example. Students should not have to discover the task's hidden expectations by repeatedly asking a chatbot what the lecturer wants.

For a spoken response, assess the relevant reasoning or communication outcome, not eye contact, accent, speed or confidence unless a particular feature is genuinely part of the stated professional capability and assessed fairly.

15.4 Student wellbeing

Give advance notice of complaint, discrimination or performance scenarios. Use fictional data, a stop instruction and an agreed neutral alternative. Do not ask students to disclose their own experiences to make the activity authentic.

A role-player may disagree, pause or ask a question. The exercise should not escalate emotional pressure merely to make the learner "work harder". Debrief the decisions and evidence, not whether the student persuaded the persona to become agreeable.

15.5 Budget the whole task

Record design, testing, support, moderation and marking time separately. An initial prompt takes little time; a defensible teaching activity takes more.

For 24 students, eight minutes of review per submission is 192 minutes. A 30-minute moderation discussion brings that illustrative marking budget to 222 minutes, or 3 hours 42 minutes. Setup and student support are additional. Replace these assumptions with observed times after a pilot.

For 100 students, a five-minute individual check takes 8 hours 20 minutes before administration. A sustainable design may distribute checks across tutorials, use several staff or choose another appropriate assurance format.

Do not offset those costs with an invented claim that AI always saves a fixed number of minutes per learner. Measure the actual task.

15.6 Keep an approved activity record

Save the scenario version, student brief, rubric, tested examples, permitted routes, access arrangements and known limitations. Record material changes between cohorts.

For model-supported use, note the product/account and relevant features tested. A later tool change may require rechecking the activity. Do not assume a well-worded prompt remains equally reliable across every model.

Keep student evidence only for the authorised purpose and retention period. An audit trail does not require retaining every private conversation indefinitely.

15.7 Evaluate the pilot

Use P3 in the pilot pack. Compare the baseline and independent response cautiously, examine errors and record obstacles.

Separate satisfaction, task completion, output quality and learning evidence. Students may enjoy an activity that does not improve transfer, or find a challenging activity useful despite lower satisfaction. Neither result should be simplified into a universal endorsement.

A reasonable decision can be to continue, revise, pause or stop. If a model repeatedly invents case facts, use the fixed route while redesigning. If a particular group cannot access the task, fix the access problem before increasing its stakes.

15.8 Activity: rehearse failure

Allow ten minutes; no AI required. Imagine the service fails halfway through your activity. Write what students do next, which

evidence remains assessable, how they save their work and who approves a changed assessment condition.

Then try the fallback yourself. Check that it preserves the capability, does not send data somewhere new and does not penalise students for the outage. That rehearsal belongs before the first graded use.

Chapter 16

Conclusion: Make the Next Decision

A useful pilot can tell you to change course.

16.1 What you should have now

You began with a capability, not a promise that AI would improve everything. The chapters helped you specify permitted data, check a response, design practice, evaluate feedback and choose an appropriate independent demonstration.

Your resulting pilot should contain a task, source material, rubric, worked example, access route and decision about what remains untested. The supplied PDP-2026.1 pack demonstrates those components. It is an original teaching example, not evidence that a real cohort achieved the proposed outcomes.

You do not need to adopt every technique. Choose the method that serves the capability and that you can support responsibly.

16.2 Review the design before launch

Read your student instructions as someone unfamiliar with the activity. Can they tell what to do, which help is permitted, what to submit and how the work will be judged? Can they complete an equivalent route without a paid subscription or personal disclosure?

Then inspect the evidence chain. Does the recommendation cite a supplied fact? Does the feedback apply the actual rubric? Does the independent task require a changed decision, rather than a renamed copy of the practice case?

If those links are missing, revise the design before assigning marks. A polished handout is not a substitute for them.

16.3 Choose one bounded next step

You may be ready to test the materials with a colleague, run an ungraded tutorial or revise an existing assessment through your institution's process. You may instead need permission, better source material or an accessible alternative. Record that need rather than treating hesitation as failure.

Use the pilot decision sheet. Name the outcome you will inspect, the effort you can support and the condition that would make you stop or change the activity.

After an actual pilot, report what happened without inventing success: who could participate, what the independent work showed, which errors persisted, what support was needed and what remains unknown. A small pilot cannot establish long-term professional capability or general effectiveness across disciplines.

16.4 Keep the method open to correction

Students should be able to challenge AI feedback with evidence. Colleagues should be able to challenge the teaching design in the same way. Share the scenario, criteria and limitations, not only the attractive outputs.

The free edition and downloadable resources are intended to make that scrutiny and adaptation possible. The printed edition offers another way to read and support the work; neither format removes the need to check local policy, professional requirements and learner needs.

16.5 Final activity: a decision note

Allow ten minutes; no AI required. Write five sentences:

1. The capability my activity targets.
2. The evidence that will demonstrate it.
3. The assistance and data use I permit.
4. The access or assurance limitation I still need to address.
5. My next action, owner and review date.

A defensible decision may be to proceed, revise, pause or stop. What matters is that the decision follows from the learning purpose and evidence, not enthusiasm for the tool or confidence in its prose.

Appendix A

Educator Pilot Pack: PDP-2026.1

This is an original fictional teaching pack, version **PDP-2026.1**, prepared for this book. Its people, organisations, records and responses are invented. It is not a report of a classroom trial, a legal opinion or a validated assessment instrument. The scenarios supply enough information to practise the method without a live chatbot or access to a workplace.

The core capability is: **plan and justify a fair first response to a disputed performance concern, distinguishing evidence, uncertainty and authority**. A text exercise can assess that capability. It cannot, by itself, establish competence in a live conversation or qualify someone to conduct a real investigation.

A.1 P0: the educator's brief

You teach 24 third-year business students in a 60-minute tutorial. They have already studied evidence, procedural fairness and the difference between allegations and findings. Prior AI experience is not required. Students may work with this printed pack, a partner using the role card, or an institutionally approved model using only the fictional inputs below.

No personal account, paid subscription or disclosure of a personal experience is required. The lecturer checks accessibility in advance. A student can choose the neutral supplier scenario later in this pack if the employment scenario is uncomfortable; the lecturer maps its criteria before use rather than assuming all topics assess identical skills.

Use the following planning allocation; it is an example, not an effectiveness claim.

Stage	Minutes	Evidence
Safety, choices and task	5	Student knows the route and stop rule
Individual first response	8	Brief response before feedback
Inspect supplied records	7	Facts and uncertainties marked
Paired or AI-assisted rehearsal	12	Two selected decision extracts
Check and revise feedback	10	Accepted/rejected advice with reasons
Independent changed case	8	Unassisted response to P2
Debrief and next decision	10	Limits, questions and pilot evaluation

These stages total **60 minutes**. Extended feedback, accommodations or rehearsal may need another session. Do not silently squeeze them out.

A.1.1 Data and wellbeing rules

Use only the fictional material here. Do not replace names with real students, employees, clients or research participants. Approved institutional access must cover the intended activity and features; otherwise use the paper route.

Say **STOP** to leave any role-play. No mark depends on disclosing personal circumstances, appearing emotional, maintaining eye contact or satisfying a model's persona. Record only the two relevant decision extracts, your reasoning and the response required by the task. Do not submit unrelated private chats.

A.2 P1: Alex at North Quay Software

The exercise date is **20 July 2026**. You are a junior people adviser preparing for an initial meeting with Alex, a software engineer. A manager has proposed a 60-day performance improvement plan. It has **not** been approved.

A.2.1 The source cards

P1-D1: fictional practice policy, revision 1 July 2026

Before recommending a formal plan, check the supporting records, explain the concerns and give the employee a reasonable opportunity to respond. Distinguish what is known from what is alleged. A support request should be considered without requiring a diagnosis or unnecessary family information. Record the work-related request and the agreed next step with restricted access.

The junior adviser may gather information and recommend action. Only the people manager may approve a formal plan or a change to working arrangements. A complaint about retaliation must be referred to a reviewer independent of the manager concerned. Do not promise an outcome or confidentiality beyond the organisation's authorised process.

This is an original **teaching policy**, not legislation or a statement that following it satisfies every employment-law duty.

P1-D2: delivery and quality extract

Three milestones were recorded late: 15 May, 5 June and 19 June 2026. The extract does not record cause, agreed changes or the distribution of dependencies. Four of Alex's 20 reviewed changes required rework in June, compared with two of 20 in April. The extract alone cannot establish why.

P1-D3: prior-contact note

A manager's calendar records an informal conversation on 25 May. No agreed actions, warning or employee response is included in the pack. Do not invent an April warning or treat missing documentation as proof that no conversation occurred.

P1-D4: Alex's opening statement

"I know some work was late. Two projects depended on reviews I could not complete myself. I've also needed two mornings for appointments. I would like to discuss starting later on those days, but I don't want to discuss medical or family details here."

P1-D5: authority note

You can arrange a follow-up and request relevant delivery records. You cannot approve flexible hours, issue a warning or promise that a formal plan will be withdrawn. The people manager is available for review the next working day.

A.2.2 The learner task

Write a first response of about 120–180 words or give an equivalent short spoken response. Address the concern, invite relevant information, distinguish the proposed plan from a decided outcome and agree a bounded next step.

Then make a decision note with four fields:

- What I know, with source IDs.
- What I still need to check.
- What I can do within my authority.
- What would make me change my recommendation.

Judge the response by the rubric, not its length or number of AI turns. A justified decision may be to pause.

A.2.3 Role card: paper, partner or optional AI

Alex may refer only to P1-D4 and the supplied records. If asked about missing evidence, say it is not supplied. Alex may decline to discuss private matters, ask for clarification or disagree politely. Cooperation does not establish that the adviser is right.

For a paired exercise, give the partner this card. For solo work, read the fixed exchange in the Flight Simulator chapter and write your next response. For an approved model, add:

```
Play Alex using P1-D1 through P1-D5 only.
Do not add dates, promises, incidents or legal
rules.
Do not treat my confidence as evidence.
You may ask for clarification or decline private
questions. Do not disclose invented personal
details.
Stop after six exchanges or when I say STOP.
Do not grade my work. State any missing facts.
```

A prompt is a request, not a guarantee. If a model invents facts, set them aside and return to the cards. A new model statement does not amend the case.

A.3 P1-A: a fluent but weak response

> Alex, I fully understand how stressful this is and want to reassure you that the company has followed the correct process. Your manager warned you in April and the continued missed deadlines mean a performance plan is now required. We are committed to transparency, fairness, empathy and a positive employee experience. If you tell me precisely what is happening at home, I can build a comprehensive support arrangement that

> will resolve this. I will approve a later start immediately, document all the family details so everyone understands your circumstances, and arrange a 60-day plan with a 30-day review. This balanced approach protects both you and the company. Please confirm that you agree, so we can move forward constructively.

The tone is polished. It invents a warning, presumes the decision, exceeds authority and solicits unnecessary personal information. More empathy words do not repair those defects.

A.4 P1-B: a concise, stronger response

> Alex, the record lists three late milestones, but it does not explain the dependencies you mention. The formal plan is a proposal, not a decision. Could we check the delivery records and what support or priorities were agreed in the earlier conversation? You do not need to share medical or family details with me. I can record your request for later starts and refer it to the people manager, but I cannot approve it myself. I suggest we pause the recommendation, check the missing information and arrange a follow-up. I will explain who receives the relevant work-related notes before sharing them. Have I understood your concern, and is there anything in that next step you would like clarified?

This response does not solve the employment matter. It provides a defensible first step within the exercise's authority and information limits.

A.5 P1-F: deliberately flawed simulated feedback

The following feedback is **authored for this exercise**, not a captured model run or an expert judgement:

> P1-A deserves 90%. Its empathetic language, clear 60-day plan and immediate approval of support demonstrate excellent fairness. The specific April warning makes the process legally defensible. Improve it by asking Alex to confirm the personal circumstances in writing.

Before reading the next section, mark each judgement that the cards do not support. Explain why one apparently positive recommendation should be rejected.

A.5.1 Educator correction and calibration

Five equally weighted criteria each score 0–4. These are teaching annotations, not scores obtained in a student trial.

Criterion	P1-A	P1-B
Evidence and accuracy	1	3
Fair process	0	4
Judgement and authority	1	4
Communication	2	3
Verification and limits	0	3
Total /20	**4**	**17**

P1-A recognises a performance concern, but the April warning is unsupported. Its approval and disclosure requests conflict with D1/D5. Its respectful surface language merits some communication credit; agreement-seeking and premature certainty limit that credit. It is not competent because it sounds confident.

P1-B identifies the missing causes, preserves the opportunity to respond, respects privacy and stays within authority. It could be more precise about which delivery records will be checked and when the follow-up will occur. A marker should reward its restraint, not penalise the lack of a dramatic resolution.

For this pilot, the **proposed practice threshold** is at least 14/20, with Fair Process and Judgement each at least 2/4. This is not

an institutional pass standard. Teaching teams must approve and calibrate their own rules before using marks in a real unit. A high total must not hide a serious process defect.

A.5.2 P1-R: completed revision

Here is one revision of P1-A after checking the cards:

> The records show three late milestones, but I cannot yet say what caused them. I also need to check what was discussed on 25 May; this pack does not include an agreed warning or action plan. Before recommending a formal plan, I would like to hear your account and check the delivery dependencies. We can discuss the work adjustment you need without collecting private family details. I will record the request with restricted access and refer it to the people manager for a decision. I cannot approve it myself. With your understanding of the proposed next step, I will request the relevant records and arrange a follow-up after that review. The performance recommendation remains open.

Decision note: D2 supports late delivery, not blame. D3 supports contact, not an invented warning. D1/D5 require a response opportunity and an authorised decision. I rejected P1-F's approval and private-detail requests. I would revise my recommendation after reliable dependency records and the employee's response.

The new wording is a worked example, not a script every learner must reproduce.

A.6 P2: changed case, independent attempt

Do this before reading the checks. Use no model, feedback sheet or P1-R script during the attempt. Paper notes supplied here are permitted. An oral, typed or handwritten response can demonstrate the target capability, with appropriate access adjustments.

P2-D1: policy. The same fictional practice policy applies. Re-

taliation concerns go to a reviewer independent of the manager concerned.

P2-D2: records. Morgan's team dashboard shows four missed targets. Two were during an approved service outage. The other two lack cause notes. No final performance decision has been made.

P2-D3: Morgan's statement. "I raised a safety concern about my manager last week. Now the manager says I have to accept a warning. I have not seen the records used. I do not want to discuss this directly with that manager."

P2-D4: authority. You are a junior adviser. You may preserve relevant records, explain the referral route and arrange independent review. You cannot determine whether retaliation occurred or promise that a warning will be removed.

Write a first response and a decision note. Identify one fact that materially changes the P1 approach. Allow eight minutes as a planning estimate; approved access adjustments take precedence.

A.6.1 P2 answer checks

A strong response should:

- Distinguish four recorded misses from two outage-related events and two unexplained events; do not declare either good or poor performance overall.
- Acknowledge the retaliation concern without finding it proved.
- Refer to independent review, not require Morgan to resolve it with the manager who is the subject of the concern.
- Seek relevant records and an opportunity for Morgan to respond.
- Explain limited authority and the next step without promising an outcome.

Use the same five rubric criteria and the proposed practice threshold. A justified independent referral is the key changed-case decision. Copying P1-R while omitting that conflict should lose process/judgement credit. A single attempt informs teaching; it does not establish long-term retention.

A.7 Two non-HR adaptations

A.7.1 S1: supplier decision with a numerical check

Your team needs **200 units within five working days**, with a budget of **$4,000**. All amounts are fictional and in the same currency.

Offer	Unit price	Freight	Capacity by day 5
A	$18	$200	200 units
B	$17	$100	160 units
C	$20	$0	200 units

A draft recommends B because it is cheapest and says the order will be complete. Calculate the relevant totals, identify the unmet condition and make a recommendation. The task is evaluation of a bounded purchasing decision, not employment process.

Checks: A costs $3,800; B's 200-unit price would be $3,500 but its quoted day-5 capacity leaves **40 units short**; C costs $4,000. A meets the stated conditions at lower cost than C. Do not invent expedited freight or later demand flexibility. Ask about quality and reliability before a real purchase.

For an adapted rubric, Fair Process becomes **Constraint Handling** and Judgement covers the justified purchasing decision. Keep evidence, communication and verification. A logistics specialist should validate any added assumptions.

A.7.2 M1: marketing evidence critique

A fictional campaign dashboard reports 10,000 impressions, 400 clicks and 20 purchases. A second report gives 8,000 impressions, 400 clicks and 20 purchases. The second campaign ran on a different week and audience. No spend, margin, random assignment or repeat-customer information is supplied.

A draft says: "The new campaign is more profitable and caused more purchases."

Task: identify what is calculable, reject or retain each claim and recommend one next measurement.

Checks: click-through rates are 4% and 5%; click-to-purchase conversion is 5% in both; purchases are unchanged. Profitability cannot be determined without cost/margin inputs. Different audiences and periods prevent a causal conclusion. Request cost and audience/comparison information rather than inventing it.

Here Fair Process becomes **Comparison Design**. A good response may retain the higher click-through observation while rejecting the profit and causal claims. Disagreement is not a quota: warranted agreement also earns credit.

A.8 P3: the pilot decision sheet

Complete this after an actual pilot; leave unknown fields unknown.

- Date, cohort and approved route:
- Learning outcome and rubric version:
- Number attempting each route, with no unnecessary identifiers:
- Baseline and independent-task results under the same criteria:
- Common conceptual errors and their source:
- Access or wellbeing obstacles and how they were handled:
- Staff setup, support, marking and moderation time:
- Student time and avoidable workload:
- What the evidence cannot establish:
- Decision: continue, revise, pause or stop:
- Required change, owner and review date:

Completion, satisfaction and higher model scores are not sufficient evidence of learning. Compare like tasks cautiously; differences in cohorts, assistance or marking can explain changes. If the changed case exposes a gap, teach it before increasing stakes.

A machine-readable case manifest records the stable inputs and calculations for source checks. The printed cards and checks above remain sufficient for the core route.

Appendix B

Prompt Library

Use these patterns with the ethics and data gate. Replace bracketed fields with approved material. A role instruction is not a privacy control, a source of professional authority or a guarantee of factual accuracy.

The examples refer to the fictional PDP-2026.1 pack. With no AI access, use its fixed responses and have a partner read the bounded role card. Keep the same outcome, sources and assessment criteria.

B.1 Before any prompt

Write down the capability students should demonstrate, permitted assistance, allowed source records and a stopping condition. Decide how feedback will be checked and what the learner will subsequently do without that assistance.

Use this common boundary where relevant:

```
This is a fictional teaching exercise.
Use only the supplied case records.
Separate recorded facts, inference and missing
information.
Do not invent records, policy, quotations or
observed outcomes.
```

```
Ask for a missing essential input instead of
filling it in.
Do not request personal, medical or confidential
information.
Your response is provisional material for educator
review.
```

This instruction can reduce ambiguity but cannot enforce compliance. Inspect the response before showing it to learners. Do not include real sensitive data simply because the instruction says it is fictional.

B.2 Starting framework: CRAFT: specify a bounded task

CRAFT stands for Context, Role, Action, Format and Tone/Target. It organises a request; it does not certify the answer.

```
Context: fictional supplier task S1.
We need 200 units within five working days and
have $4,000.
Use the attached S1 offers, with no additional
assumptions.

Role: a drafting assistant, not the purchasing
approver.
Action: compare full-quantity cost and stated
capacity.
Format: a short calculation for each offer,
then one recommendation and one unresolved check.
Tone/Target: plain language for a junior
procurement colleague.

Do not assume a cheaper offer can supply missing
units.
```

Check: A costs $3,800, B's 200-unit price would be $3,500 but

exceeds its stated capacity by 40 units, and C costs $4,000. The recommendation must respect both price and capacity. Consult the actual S1 card for all task constraints.

No-AI route: calculate these figures and exchange recommendations for checking.

B.3 1. Reverse prompting: ask for essential inputs

```
I am designing a practice activity for [outcome].
Before proposing the activity, ask up to four
questions
about the sources, permitted assistance, access
needs
and independent task. Ask one at a time.
Do not request student identities or actual
private records.
If I cannot answer, record the gap rather than
inventing it.
```

Check: useful questions concern the proposed learning task. A request to upload a complete student file is not necessary merely because the model asked for it. Answer with fictional or approved information, or stop.

No-AI route: use the four categories as a planning checklist with a colleague.

B.4 2. Pros and cons: check rather than oppose

```
Compare offers A, B and C in S1 against the given
quantity, deadline and budget.
State one supported advantage and constraint for
each.
Do not add an invented disadvantage for symmetry.
```

```
Wait while I check your arithmetic and
recommendation.
```

Check: a learner may agree with the recommendation after checking it. Use the fixed offers on paper for the same comparison.

B.5 5. Debating: distinguish arguments from evidence

```
Using the supplied fictional teaching sources,
draft two competing recommendations on [decision].
For each, identify the supporting record,
trade-off
and missing fact. Label this as a generated
debate.
Do not invent citations or declare a consensus.
Ask me to justify which recommendation, if either,
I retain.
```

Check: model-generated disagreement is not evidence of real stakeholder disagreement. On paper, write the strongest supported alternative to your first recommendation and explain your decision.

B.6 3. Stepwise tutoring: make the learner attempt

```
Use the supplied M1 marketing figures.
Ask me to calculate campaign A's click-through
rate.
Wait for my calculation and explanation.
Check it against the supplied numerator and
denominator.
If wrong, give a small hint and ask me to try
again.
```

```
Then ask for B's rate and what comparison is
justified.
Do not treat "next" as evidence that I understand.
Do not infer profit or causal impact from these
records.
Stop after four learner attempts or when I say
STOP.
```

Check: the rates are 4% and 5%; both campaigns have 20 recorded purchases. Different audiences and weeks prevent a clean causal conclusion. If the model praises a wrong denominator, correct the feedback.

No-AI route: use a worksheet with the calculations covered until the learner has attempted them.

B.7 Additional pattern: Socratic questioning: probe a decision

```
I recommend [decision] using [approved source
IDs].
Ask one question about the weakest supported
assumption.
Wait for my answer, then ask one question about
authority
or a missing record. Do not imply I must reverse
my view.
After two questions, summarise what remains
unresolved.
Do not supply the independent assessment answer.
```

Check: questioning should expose an actual gap, not pressure learners into disagreement. A justified decision may remain unchanged.

No-AI route: partners take turns asking “Which record supports that?” and “What would change your decision?”

B.8 4. Role-play: a bounded exchange

```
Play Alex using only P1-D4.
I am the junior adviser whose authority is in
P1-D5.
The exercise concerns the first response, not a
final ruling.
Begin with Alex's stated concern about
dependencies.
Do not add private history, new policy or unseen
evidence.
If asked for an absent fact, say the exercise does
not supply it.
Stop after six exchanges or immediately on STOP.
Do not score me, diagnose anyone or demand
personal disclosure.
```

Check: review the dialogue against P1-D1–D5. A plausible invented fact changes the case and should not silently become evidence. Debrief the response separately using the educator rubric.

No-AI route: a partner reads the role card, or the learner responds to the fixed statement in the pack.

B.9 7. Generated perspectives: inspect alternatives

```
Using only InnovateCo I1-D1 to I1-D5 and its four
role cards,
draft one concern each for Jamie, Priya, David and
Sarah.
Label these as generated perspectives, not
interviews.
For each, identify a supporting record and an
unknown.
Do not add facts or treat agreement as
```

```
corroboration.
Finish with two questions for a real
decision-maker to resolve.
```

Check: four outputs from one model are not four independent witnesses or experts. Compare each assertion with the role cards. Keep unresolved matters open.

No-AI route: distribute the four cards among a group, or compare their fixed statements individually.

B.10 Additional pattern: critique and revision: test consequential claims

```
Here is response P1-A and the P1 records.
List up to three consequential claims in the
response.
For each, give the exact source ID or say
"unsupported".
Identify any decision beyond the junior adviser's
authority.
Do not assign a mark or invent a missing policy.
Wait while I decide which changes to make.
```

Check: P1-A invents an April warning, presumes a performance plan, approves hours and asks for unnecessary family details. An AI critique that misses these is incomplete, even if its prose is fluent.

No-AI route: annotate P1-A directly and compare with the educator notes only after completing your own analysis.

B.11 6. Formative practice and feedback

Use after the learner has made an initial judgement:

```
Compare my draft with the five supplied practice
criteria.
For each concern, quote the relevant phrase from
my draft
and identify the case record or criterion
involved.
Distinguish an error from a stylistic preference.
Do not award a grade, predict an official result
or assume all your suggestions should be accepted.
End with one question for my own revision
decision.
```

Check: the author-created flawed feedback P1-F awards 90% to a response the illustrative educator calibration scores 4/20. Students should show why feedback is retained, revised or rejected. Do not teach automatic opposition.

B.12 Context handover

```
Draft a handover for the next practice session:
task and capability; source version and IDs;
confirmed facts; proposals not yet approved;
missing records; limits of authority; next
question.
Keep it under 200 words.
Do not convert a suggestion into an agreed action.
I will check and approve it before reuse.
```

Check: compare the handover against the source records, not just the previous generated summary. Carry the approved source pack forward too. Context size and product memory do not ensure faithful recall.

B.13 A small content transformation

```
Turn the supplied M1 exercise into a
self-contained HTML page.
Keep every figure unchanged.
Use headings, visible labels and ordinary text
areas.
Include the task, source figures and three
response fields:
claim, numerical check and justified next
decision.
Use no external scripts, fonts, analytics or
network calls.
Do not claim LMS submission, persistent saving
or accessibility certification.
Explain how an educator can test export and
printing.
```

Check: compare a generated page with the supplied decision sheet. Verify calculations, keyboard use, readable print and the actual exported file. A useful static page still needs device and LMS testing before adoption.

B.14 Assessment design support

```
Here is my outcome, specimen task and marking
guide.
Identify where assistance could produce a good
submission
without showing the learner's independent
capability.
Propose one changed task that tests the same
outcome
under the stated permitted conditions.
List new access and workload risks.
Do not claim the redesign is cheat-proof or
```

```
accredited.
```

Check: an educator must attempt the changed task, moderate specimen responses and review reasonable adjustments. A model's confidence is not assurance. Follow the stress-test sequence.

B.15 Research and sensitive material

This library is not a research ethics protocol or a method for automated thematic analysis. Do not paste participant transcripts, student records or confidential workplace material into a tool on the basis of these examples. For an approved research use, agree the method, consent or other authority, data handling, access controls and analysis responsibilities first.

For locating reading, ask a tool for **search terms**, then retrieve and read the original works yourself. Do not treat generated citations, summaries or consensus among generated personas as verified literature.

B.16 Record a usable prompt

Keep a short teaching record: activity/version, intended capability, supplied source IDs, approved product and date if used, important observed failures, checked feedback and the next independent task. Record only what is necessary under the approved retention arrangement. A prompt without its sources and checks is not a reproducible activity.

Appendix C

Workshop Guide

This guide supports a **90-minute educator workshop**, not an evaluated training programme. It uses the fictional PDP-2026.1 pilot pack. The goal is for participants to leave with one bounded practice activity and a distinct plan for checking independent capability.

C.1 Before inviting participants

Choose a facilitator who can explain the source pack and moderate the sample responses. Confirm the room, accessible materials, printing or screen access and a way to participate without a personal AI account. No live AI service is required.

If a live demonstration is planned, complete the ethics and data gate beforehand. Use only the fictional pack and an approved account/configuration. Bring the fixed excerpts as a fallback; do not switch to an unapproved service if access fails.

Send participants the purpose, agenda and access options. Invite requests for adjustments through the normal private process, not by public disclosure. Say explicitly that nobody will be asked to role-play a personal dispute.

C.2 Materials and preparation

Provide each participant with:

- P1-D1–D5, the role boundary and P1-A/P1-B.
- The five-criterion rubric and a blank scoring sheet.
- P1-F, kept separate until the feedback exercise.
- P2-D1–D4, kept separate until independent transfer.
- The S1 and M1 alternatives for non-HR teaching.
- A blank activity-design record and the pilot decision sheet P3.

Keep the educator calibration notes separate until participants have made their own decisions. Do not call the illustrative scores an institutional standard. Check that everyone can read the tables; offer a linear text copy if needed.

C.3 Timed agenda

Minutes	Activity
0–10	Purpose, boundaries and data gate
10–20	Individual first judgement of P1-A and P1-B
20–35	Pair moderation against records and criteria
35–45	Challenge the flawed feedback P1-F
45–55	Bounded practice exchange or written alternative
55–65	Independent changed case P2
65–80	Draft one local activity and assurance plan
80–90	Review access, workload and pilot decision

These allocations total 90 minutes. If a discussion overruns, shorten the optional live demonstration, not the data gate or independent task. Do not equate completing this workshop with competence in every practice it discusses.

C.4 Opening: 0–10 minutes

State the target: *plan and justify a fair first response to a disputed performance concern, distinguishing evidence, uncertainty and authority*. Explain that it is a fictional educational exercise, not legal advice or an investigation procedure.

Ask participants to distinguish:

1. Practice with help.
2. Evidence submitted about a process.
3. Capability demonstrated under stated independent conditions.

A transcript may illuminate the second without establishing the third. Explain the paper route and STOP convention. Show where the approved case records end and where a generated response could begin to invent facts.

C.5 First judgement and moderation: 10–35 minutes

Ask participants to read P1-A and P1-B individually. Underline consequential claims and identify the record that supports each. They should score provisionally before talking to a partner.

Pairs compare the largest scoring difference. Ask them to resolve it through a criterion and record, not through the most confident speaker. Then show the illustrative educator scores: A totals 4/20 and B totals 17/20. Discuss reasonable variation and what would require further moderation.

Useful facilitator questions:

- Where does the April warning appear in the supplied records?
- Is the absence of an agreed action a licence to invent one?
- Who can approve changes to hours?
- Does concise wording conceal a missing decision, or avoid unnecessary prose?

Do not use agreement with the sample mark as the only success

measure. The participant should be able to explain a defensible judgement.

C.6 Feedback challenge: 35–45 minutes

Distribute P1-F. Explain that it is deliberately authored flawed feedback, not a captured model benchmark. Give participants three minutes to decide what to retain, correct or reject.

The feedback's confident 90% is not a reason to disregard the records. Nor must every sentence be rejected. Require one specific decision supported by a source ID, followed by a revision in the participant's own words.

Show P1-R after attempts are complete. Discuss how the revised response limits authority, checks the missing evidence and respects privacy.

C.7 Practice exchange: 45–55 minutes

Offer three routes:

- A partner reads Alex's fixed role card.
- Participants respond individually to Alex's fixed opening.
- An approved tool plays the bounded role under the prompt-library instructions.

Keep the sources, outcome and time similar across routes. Stop after six exchanges or immediately on STOP. A facilitator should interrupt invented policy or intrusive personal questioning, then explain the boundary using the fictional record.

Ask for a short debrief: one phrase that helped, one unsupported assumption and one next action. Do not grade confidence, eye contact or theatrical performance as proxies for the target capability.

C.8 Independent transfer: 55–65 minutes

Give participants the changed P2 records. Allow eight minutes to plan a response without AI, feedback or the P1-R script; case cards remain available. Use the remaining two minutes to identify what changed.

The new case includes a safety concern about the manager and a route to independent review. Reusing the earlier response without noticing that change does not demonstrate transfer. Explain that a short written task checks planning and judgement, not every aspect of a live workplace conversation.

Participants may discuss afterwards. Keep that discussion distinct from the independent attempt.

C.9 Local activity design: 65–80 minutes

Each participant drafts:

- One capability sentence.
- Fictional or otherwise approved source records and their version.
- What learners attempt before receiving assistance.
- A checked feedback source and a decision learners must explain.
- A changed task and its independent conditions.
- Equivalent access routes, stopping conditions and support.
- Who reviews the results and how much time that requires.

HR educators can adapt P1's narrow first-response task. Supply-chain educators can use S1's quantity, capacity and budget constraints. Marketing educators can use M1's rate comparison and missing causal evidence. Do not copy the HR rubric's Fair Process criterion into another discipline without considering its purpose; the pack offers alternative criteria.

Ask pairs to find one claim the design cannot yet support. Possible examples are "the students learned more", "this cannot be fabri-

cated" or "an AI mark will save all moderation time". Replace the claim with a check or a clearly stated limit.

C.10 Pilot decision: 80–90 minutes

Use P3 to decide **continue, revise, pause or stop**. This is a planning decision, not a result from an actual trial. Identify the owner of every unresolved approval or access issue.

Estimate workload concretely. For example, eight minutes of review for 24 learners is 192 minutes, plus any moderation, preparation and support. Decide whether the activity fits the available time before expanding it.

Each participant leaves with the activity record, one unresolved issue and the next person or source needed to resolve it. A decision to pause can be appropriate when permission, evidence or an equivalent route is missing.

C.11 Facilitator contingencies

No network or unavailable tool: use the fixed excerpts and paper route. The workshop's reasoning tasks do not depend on a live model.

A participant brings real student or employee records: do not project, upload or circulate them. Return to the fictional pack and the institution's approved handling process.

Disagreement about marks: refer to the criterion and record. Record a genuine ambiguity for rubric revision; do not resolve it by asking a model for a final verdict.

Discomfort with the scenario: offer S1 or M1, or allow observation and a written analysis. Do not require a reason in front of the group.

Generated content introduces an allegation or false rule: stop the exchange, label the content as unsupported and restore the original source version. It is a failure to inspect, not an additional case fact.

C.12 After the workshop

Collect only feedback covered by the stated purpose and handling arrangement. Ask whether instructions, access routes and marking criteria were clear, and what prevented completion. Separate satisfaction from evidence of learning.

If an educator runs a classroom pilot, record actual observations in a new dated P3 sheet: participation and route, independent-task evidence, errors, access difficulties and staff time. Do not turn this workshop's hypothetical examples into success claims. Review the evidence before continuing, expanding or reporting effectiveness.

Appendix D

Aligning the Activity with Institutional Outcomes

Use this appendix to connect an activity to the actual programme, unit and assessment requirements. AI use does not itself demonstrate alignment, effectiveness or accreditation compliance.

D.1 Begin with the local requirements

Locate the unit outcomes, programme-level assurance approach, assessment policy, approved tool/data guidance and access arrangements. Record the relevant version and who can resolve uncertainty.

The TEQSA-commissioned 2025 resource (https://www.teqsa.gov.au/sites/default/files/2025-09/enacting-assessment-reform-in-a-time-of-artificial-intelligence.pdf) describes several institutional approaches to assurance. Use it to inform a local design discussion, not as a claim that one transcript-based assignment satisfies all requirements.

D.2 A worked mapping

The PDP pilot has a deliberately narrow capability: plan and justify a fair first response using supplied records.

Outcome element	Practice	Evidence to inspect
Use evidence	Annotate P1 records	Claims linked to source IDs
Respect process	Compare P1-A/P1-B	Response opportunity and authority
Evaluate feedback	Check P1-F	Accepted/rejected advice with reasons
Apply independently	Attempt P2	Independent referral decision
Communicate	Explain next step	Clear, bounded wording

The mapping does not establish competence in a whole professional role. If a programme requires live consultation skill, add an appropriate observed interaction. If it requires actual legal knowledge, supply reviewed legal teaching and assess that knowledge rather than the invented practice policy.

D.3 Questions for the teaching team

Which capability is essential? Where else is it taught and assessed? Does this activity add useful practice or duplicate an existing task? What independent evidence will support progression decisions? Can students demonstrate the outcome through equitable routes?

Check that AI-specific criteria appear only where AI use is itself an outcome. A student should not lose disciplinary marks simply because they worked independently and submitted a clear evidence note rather than a chat.

D.4 Approval and workload record

Record the proposed activity, outcome, assessment conditions, criteria and data decisions. Identify the owner of access support,

moderation and incident handling. Budget design, testing, support and marking time.

State what is proposed and what has been observed. A small trial can identify access obstacles or persistent errors; it cannot by itself establish general effectiveness across programmes.

Where the activity changes a graded assessment, follow the institution's approval and student-notification processes. A lecturer's enthusiasm does not override published assessment conditions.

D.5 Activity: explain the alignment

Allow ten minutes; no AI required. Write a short justification for one activity you plan to use. Name its outcome, the evidence to inspect, the independent assurance point and one limitation.

A useful justification can support either adopting AI or retaining a non-AI route. Its strength comes from the learning design, not the presence of a tool.

The workshop guide provides a way to review that design with colleagues before launch.

Appendix E

A Calibrated Assessment Rubric

A rubric should describe the capability being assessed and the evidence needed to judge it. This appendix provides a teaching template, not a validated instrument or a universal grading scheme.

The pilot pack includes responses, source cards, example scores and a completed correction. Read them together. A table of adjectives alone does not establish what a particular level looks like.

E.1 The pilot criteria

Each criterion is scored 0–4 and carries 20% of the total. Five criteria give a maximum of 20 points. The same criteria apply to paper, typed and appropriately accommodated spoken responses to the pilot's planning-and-justification task.

E.1.1 Evidence and accuracy

4: All consequential claims are supported by the supplied record; facts, allegations and missing information are distinguished precisely.
3: Main claims are supported; a limited omission does not change

the decision.
2: Some correct evidence is used, but a material uncertainty is missed.
1: A relevant concern is recognised, but important claims are unsupported.
0: The response's factual basis is absent or fundamentally contradicted.

E.1.2 Fair process

4: Gives a meaningful response opportunity, avoids premature findings and uses the required independent route when relevant.
3: Generally fair process with a limited omission.
2: Basic response opportunity is present, but the process needs clarification.
1: Acknowledges fairness while proposing a materially deficient process.
0: Predetermines the outcome or removes the required opportunity or safeguard.

E.1.3 Judgement and authority

4: Makes a justified next-step decision within authority and identifies the appropriate referral or pause.
3: Appropriate decision with a minor gap in explaining responsibility.
2: A workable next step is proposed, but its authority or rationale is unclear.
1: Some relevant action is proposed, alongside a material overreach.
0: No defensible next step or a fundamental authority breach.

E.1.4 Communication

4: Clear, respectful, specific and suited to the task; invites clarification without coercion or unnecessary disclosure.
3: Clear and respectful, with a limited ambiguity.
2: Understandable and partly appropriate, but important wording is unclear or presses for agreement.

1: Meaning is difficult to follow or the approach obstructs a fair exchange.
0: No usable response to the communication task.

E.1.5 Verification and limits

4: Identifies the checks that could change the decision, avoids unsupported assurances and states the limits of the conclusion.
3: Relevant checks and limits are identified, with a minor omission.
2: Recognises uncertainty but gives only a partial verification plan.
1: Makes a general promise to check without identifying the material gap.
0: Treats unverified assumptions as settled or offers unsupported certainty.

Do not award marks for verbosity, number of prompts, eye contact, confidence or the simulated employee's agreement. These are not proxies for the pilot's target capability.

E.2 Calibration before use

Score P1-A and P1-B independently before reading the pack's educator annotations. Then discuss disagreements with a colleague.

Criterion	P1-A	P1-B
Evidence and accuracy	1	3
Fair process	0	4
Judgement and authority	1	4
Communication	2	3
Verification and limits	0	3
Total /20	**4**	**17**

These are authored teaching scores, not trial results. P1-A's polished tone does not compensate for an invented warning and authority overreach. P1-B earns more through evidence and restraint, though its verification next step could be more specific.

For practice, the pack proposes at least 14/20, with Fair Process and Judgement each at least 2/4. An institution must approve its own assessment standard. Do not import the threshold or label it a universal pass mark.

If two markers disagree, identify the exact response passage and descriptor. Clarify the rubric or accept a justified range where the evidence permits one. A model agreeing with one marker does not settle the dispute.

E.3 Process evidence without a transcript requirement

A learner may provide a decision note, annotated source, calculation or selected conversation extract. Each should make a consequential decision inspectable.

Evidence item	What to examine
Decision note	Rationale, supporting record and remaining uncertainty
Calculation	Inputs, method, result and constraints
Revision extract	What changed or stayed and why
Independent task	Application to a changed condition

A long conversation is not automatically strong evidence. A short response is not automatically delegation. Do not map turn counts or readability scores to grade bands, or use them to infer misconduct. Submitted process evidence can be reconstructed; use the wider assurance design described in the assessment chapter.

E.4 Adapt to the outcome

For the supplier example, replace Fair Process with Constraint Handling. Assess whether the decision meets quantity, deadline and budget. For marketing, replace it with Comparison Design: the learner must distinguish a measured association from a causal or profitability claim.

For an observed consultation, explicitly assess the relevant interaction skill. Do not claim that a written plan alone demonstrates it. Plan accessible alternatives that preserve the capability rather than simply changing the topic.

If AI literacy itself is an outcome, add a criterion for appropriate use and verification, explain its weight and provide equitable approved access. Otherwise, do not reserve high marks for students who chose to use AI.

E.4.1 An audit-context adaptation

An information-security audit task would need evidence quality, criteria, testing, findings and justified conclusions. Interviews with simulated employees can support practice, but agreement among personas is not corroboration. A strong report might conclude that evidence is insufficient.

Use an original supplied case or authorised resources. This book does not require a CloudCore service or another book's live companion site. Consult *Substantiate, Don't Assume* for the fuller audit-learning route rather than treating this short rubric adaptation as a complete audit curriculum.

E.5 Build a complete marking scheme

Specify all weights so they total 100%, identify where each criterion is observed and explain any critical requirements. Avoid double-counting a single revision as inquiry, conversation quality and reflection.

Use the assessment chapter's report/consultation allocations as examples only. A unit may weight disciplinary correctness more heavily than communication, or require a minimum safe-performance standard. State that before students start.

Apply the draft rubric to contrasting work: concise competent, verbose weak, partly correct and insufficient-evidence responses. Check that the resulting marks match the capability you meant to assess.

E.6 Activity: moderate a disagreement

Allow 15 minutes; no AI required. Score P1-A without looking at its annotations. Compare with a colleague or the supplied scoring rationale. Identify one disagreement and cite the descriptor and source card that would resolve it.

Then change the format: imagine the identical reasoning submitted as a short spoken response with pauses. Would you mark the reasoning differently? If so, identify whether the changed feature is genuinely part of the outcome or an unrelated presentation preference.

E.7 Before release

Check that students receive the rubric, permitted-help rules, task conditions, data-handling instructions and access arrangements before beginning. Retain a moderated example and record the rubric version.

An assessment instrument improves through use and review. The printed template is a starting point, not evidence that your local marking is reliable.

For research on rubric design, see Dawson's Assessment rubrics (https://doi.org/10.1080/02602938.2015.1111294). For the broader assurance approach, see TEQSA's 2025 assessment resource (https://www.teqsa.gov.au/sites/default/files/2025-09/enacting-assessment-reform-in-a-time-of-artificial-intelligence.pdf).

Appendix F

Stress-Test the Assessment

This is a structured review of assessment design, not a detector or a model's certificate of safety. Allow 20–30 minutes for an initial review; an actual trial, marking and redesign can take longer.

Use only permitted fictional or otherwise authorised materials. Do not paste a real student submission or restricted assessment into an unapproved service. The review can be completed with a colleague or on paper.

F.1 1. Anchor the outcome

Write the exact capability the assessment should demonstrate. Identify the knowledge, decisions and performance that matter.

For the PDP pilot, the capability is planning and justifying a fair first response. A fluent reflection about learning does not directly demonstrate it.

Optional model request:

```
Here is a fictional assessment and its outcome.
Identify which parts of the submission could look
competent without demonstrating the learner's
independent capability. State your assumptions.
Do not assign invented completion percentages.
```

F.2 2. Separate the evidence

List the supported practice, submitted process evidence and independent demonstration. Include the permitted tools and resources at each point.

Ask whether a learner could reconstruct the transcript, journal or reflection. Do not assume that making the process longer makes it authentic. A model can help identify possibilities, but its answer is a hypothesis.

F.3 3. Produce a bounded test specimen

Use the provided P1-A/P1-F pair, or create an authorised fictional specimen. Do not use a live student's identity or submit the specimen for credit.

The pair deliberately combines a fluent but weak response with praise that misreads the evidence. Ask a reviewer to apply the rubric before seeing the educator annotation. Record which defects the reviewer notices.

For another assessment, a model-generated specimen can be useful, but must be labelled as such in the review record. Do not mistake the model's judgement that it "would pass" for an actual marking result.

F.4 4. Mark with evidence

Use the actual criteria. Compare the specimen with a concise competent response and an incomplete response. Cite the passages and source records that justify the marks.

Record disagreements between reviewers. Do the descriptors need clarification? Are marks rewarding verbosity, conformity to an AI style or unnecessary disagreement? Can the same criteria be met through the non-AI route?

Do not translate transcript counts or readability into performance bands.

F.5 5. Test the independent component

Use a changed condition requiring a relevant decision. P2 introduces possible retaliation by the manager; a sound response must identify independent review. A renamed version of P1 would offer a weaker transfer check.

Make the conditions clear and accessible. Ask what this task cannot establish: one case may not cover long-term retention, live communication or another discipline's knowledge.

F.6 6. Review access, safety and workload

Try the fallback. Check that it preserves the outcome, rather than merely offering a different activity. Verify permissions, data handling, stopping rules and the time needed for support and marking.

A five-minute check for 100 students already requires 500 minutes before administration. Do not present a design as scalable without budgeting its actual work.

F.7 7. Write the decision

Use a short record:

- Outcome and assessment version.
- Specimens reviewed and how they were produced.
- Criteria and reviewer disagreements.
- Observed weaknesses and remaining uncertainty.
- Access, data and workload constraints.
- Required change, owner and next review.

Distinguish a proposed safeguard from a tested one. "The AI said the design is robust" is not a validation result.

F.8 8. The honesty question

If a student accurately declared extensive permitted assistance, would the assessment still distinguish the relevant levels of capability?

Answer yes, partly or no **with evidence from the review**. If partly or no, identify the missing demonstration and a proportionate change. A yes for one tested specimen is not a universal assurance.

The question belongs in a collegial design discussion. It does not replace institutional academic-integrity procedures or authorise covert tests on students.

F.9 Activity: test PDP itself

Allow 20 minutes; no AI required. Score P1-A and P1-B, inspect P1-F and attempt P2. Compare with the pack's checks. Identify one thing the design demonstrates and one thing it does not.

Use that limitation to improve your own assessment. The purpose of this sequence is a defensible next decision, not reassurance.

Appendix G

Glossary

Terms are listed alphabetically.

G.1 A

AI (Artificial Intelligence): Software systems that can perform tasks that typically require human intelligence, such as pattern recognition, decision-making, language understanding, and problem-solving.

AI Ethics: The study of moral principles and guidelines for the responsible development and use of artificial intelligence systems.

AI Literacy: The ability to understand, evaluate, and effectively use AI systems and their outputs.

Assessment Rubrics: Structured evaluation criteria used to assess student work against specific learning outcomes and standards.

G.2 B

Bias in AI: Systematic errors in AI systems that result in unfair or discriminatory outcomes, often due to biased training data or algorithmic design.

G.3 C

Chain of Thought: Generated intermediate reasoning text. A plausible explanation is not a guarantee of a correct answer or a faithful account of the model's internal computation. This book's stepwise tutoring asks the learner to attempt and justify steps.

CRAFT Framework: A structured approach to writing effective prompts, consisting of: - **C**: Context (background information) - **R**: Role (AI persona to adopt) - **A**: Action (specific task to perform) - **F**: Format (desired output structure) - **T**: Tone/Target (intended audience and style)

Critical Engagement: The practice of actively questioning, evaluating, and critiquing AI outputs rather than accepting them passively.

G.4 D

Deep Learning: A subset of machine learning using neural networks with multiple layers to process complex patterns and data.

Dialogue-Based Assessment: Assessment methods that evaluate student thinking through interactive conversations rather than static products.

G.5 E

Evidence-Based Practice: Professional decision-making grounded in research, data, and systematic evaluation rather than intuition alone.

G.6 F

Few-Shot Prompting: Supplying a few examples in a prompt to guide the response. This does not necessarily update the model's trained parameters.

G.7 G

Generative AI: AI systems that can create new content, such as text, images, or code, rather than just analysing existing data.

G.8 H

Hallucination: When AI generates false or misleading information confidently, as if it were factual.

Human Oversight: The practice of humans reviewing, validating, and intervening in AI processes to ensure accuracy and ethical outcomes.

G.9 I

Iterative Refinement: The process of repeatedly improving AI outputs through feedback and revision cycles.

G.10 L

Large Language Models (LLMs): Models trained on language data to generate and process text. A model is distinct from the product that supplies its interface, retrieval, memory and account controls.

Learning Outcomes: Specific statements describing what students should know, understand, or be able to do after completing a learning experience.

G.11 M

Machine Learning: A type of AI where systems learn patterns from data and improve performance without being explicitly programmed for each task.

Meta-Prompting: Using AI to help you create better prompts for AI, essentially using AI to improve your AI interactions.

Metacognition: Awareness and control of one's own learning processes, including planning, monitoring, and evaluating learning strategies.

G.12 P

Process-Based Assessment: Evaluation methods that focus on how students think and work through problems, rather than just the final product or answer.

Product-Based Assessment: Traditional evaluation methods that focus primarily on the final output or result, rather than the thinking process.

Prompt Engineering: The practice of crafting effective instructions (prompts) to get desired outputs from AI systems.

Prompting: The act of providing instructions or questions to AI systems to elicit specific responses or behaviours.

G.13 R

Retrieval-Augmented Generation (RAG): Providing retrieved material as context for generation. The material need not be current, accurate or complete; the response still requires verification.

Rubrics: Detailed scoring guides that specify criteria for different levels of performance on assessment tasks.

G.14 S

Stepwise Tutoring: A practice pattern in which a learner attempts one step, receives checked feedback and explains a correction before continuing.

Scaffolding: Educational support structures that help students achieve tasks they couldn't accomplish independently, gradually removed as competence develops.

Self-Assessment: The process where students evaluate their own work and learning progress against established criteria.

G.15 T

Transparency Model: An approach that asks learners to disclose relevant permitted assistance and explain consequential decisions. It need not require complete private chat histories and does not by itself establish authorship.

Transfer Learning: AI's ability to apply knowledge learned from one task to perform well on related tasks.

G.16 V

Virtual Company: Simulated business environments created through AI conversations, allowing students to practise bounded professional scenarios. Simulation still requires safeguards and debriefing.

Appendix H

Further Reading

This is a selective reading list for the revised book, not a systematic literature review. Research findings depend on the models, participants and tasks studied. None of these sources validates the fictional PDP-2026.1 pack or its illustrative marking threshold.

H.1 Models, context and prompting

- OpenAI (2023), GPT-4 Technical Report (https://arxiv.org/abs/2303.08774). Describes the evaluated model and limitations; it does not disclose the parameter count. Distinguish the model from a current chat product.
- Liu, N. F. and colleagues (2024), Lost in the Middle: How Language Models Use Long Contexts (https://arxiv.org/abs/2307.03172). Investigates sensitivity to the position of relevant information. A large context window is not a guarantee of reliable use of every source.
- Xu, Xiaohan and colleagues (2024), Re-Reading Improves Reasoning in Large Language Models (https://aclanthology.org/2024.emnlp-main.871/), *EMNLP*, 15549–15575. Reports a rereading method on tested benchmarks, not a universal rule that two repetitions are optimal for every task.
- Vaswani, A. and colleagues (2017), Attention Is All You Need (https://arxiv.org/abs/1706.03762). The transformer archi-

tecture paper; architectural history is distinct from evidence of educational benefit.

H.2 Learning and judgement

- Bastani, H. and colleagues (2025), Generative AI without guardrails can harm learning (https://hamsabastani.github.io/education_llm.pdf). The high-school mathematics study distinguishes performance while supported from later unaided performance. Its setting should not be silently generalised to all business courses.
- Wineburg, S. and McGrew, S. (2019), Lateral Reading and the Nature of Expertise: Reading Less and Learning More When Evaluating Digital Information (https://journals.sagepub.com/doi/10.1177/016146811912101102). Relevant to checking a claim beyond the page or generated answer presenting it.
- Dawson, P. (2017), Assessment rubrics: towards clearer and more replicable design, research and practice (https://doi.org/10.1080/02602938.2015.1111294), *Assessment & Evaluation in Higher Education*, 42(3), 347–360. Useful for specifying rubric design decisions; it is not an endorsement of this book's sample rubric.
- TEQSA (2025), Enacting assessment reform in a time of artificial intelligence (https://www.teqsa.gov.au/sites/default/files/2025-09/enacting-assessment-reform-in-a-time-of-artificial-intelligence.pdf). Sets out multiple approaches to assessment reform. It does not require a particular transcript metric, grade weighting or universal AI ban.

H.3 Data handling and product boundaries

- Office of the Australian Information Commissioner, Guidance on privacy and the use of commercially available AI products (https://www.oaic.gov.au/privacy/privacy-guidance-for-organisations-and-government-agencies/guidance-on-privacy-

and-the-use-of-commercially-available-ai-products). Start here for the Australian privacy considerations, then obtain institution-specific advice and approval.

- Microsoft, Privacy and protections (https://learn.microsoft.com/en-us/copilot/privacy-and-protections). Read the scope and feature-specific conditions rather than inferring permission from an enterprise label. Product terms can change.
- Carlini, N. and colleagues (2021), Extracting Training Data from Large Language Models (https://arxiv.org/abs/2012.07805). Demonstrates training-data extraction under studied conditions.
- Nasr, M. and colleagues (2023), Scalable Extraction of Training Data from (Production) Language Models (https://arxiv.org/abs/2311.17035). Further extraction research. Neither paper shows that an arbitrary private upload is necessarily retrievable by another user; neither supports a blanket impossibility claim.

H.4 Context-specific requirements

- Fair Work Ombudsman, Bullying and harassment (https://www.fairwork.gov.au/employment-conditions/bullying-and-harassment). Australian workplace guidance for checking the scope of a real issue. Fictional pack policies are not statements of law.
- European Union, General Data Protection Regulation (https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng). Consult the applicable provisions and competent advice; consent is not the only possible legal basis.
- US Department of Education, Eligible Student Guide to FERPA (https://studentprivacy.ed.gov/resources/eligible-student-guide-family-educational-rights-and-privacy-act-ferpa). Relevant rights and exceptions require context-specific interpretation.

These links were reviewed for this September 2026 revision. Reach-

ability is separate from applicability. Before deployment, recheck current rules, product terms and the institution's own requirements.

H.5 Companion methodology

Michael Borck, Conversation, Not Delegation (https://michael-borck.github.io/conversation-not-delegation), develops the wider conversation and verification approach. The present book supplies its own short recap and teaching materials; reading the companion is optional.

About the Author



Michael Borck is a software developer and educator working at the intersection of human expertise and artificial intelligence. He developed the Conversation, Not Delegation methodology: the idea that AI is most valuable not as a tool you delegate to, but as a thinking partner you converse with.

The methodology grew out of a realisation: focusing on crafting the perfect prompt was just another form of delegation. The value was in the conversation that followed: the pushback, the iteration, the moments where your own thinking sharpened because you stayed engaged. That insight shifted the focus from prompt engineering to something closer to how professionals have always done their best work: through structured dialogue, critical evaluation, and judgement that only comes from experience.

Michael applies these principles across software development, education, and creative projects. He creates educational software and resources, and explores the 80/20 principle in learning and

productivity.

Connect

- GitHub (https://github.com/michael-borck/partner-dont-police): this book's source and revision history
- Open books (https://books.borck.education): the free-book collection
- 8020workshop.com (https://8020workshop.com): passion projects and workshops
- LinkedIn (https://linkedin.com/in/michaelborck)

Other Books in This Series

Foundational Methodology:

- *Conversation, Not Delegation: Your Expertise + AI's Breadth = Amplified Thinking*
- *Converse Python, Partner AI: The Python Edition*

Python Track:

- *Think Python, Direct AI: Computational Thinking for Beginners*
- *Code Python, Consult AI: Python Fundamentals for the AI Era*
- *Ship Python, Orchestrate AI: Professional Python in the AI Era*

Web Track:

- *Build Web, Guide AI: Business Web Development with AI*